

# Chapter 7: Reproducible data analysis

Mark Andrews

## Contents

<b>Introduction</b>	<b>1</b>
<b>Using RMarkdown for reproducible reports</b>	<b>3</b>
Installation . . . . .	4
A minimal RMarkdown example . . . . .	4
An extended RMarkdown example . . . . .	5
An additional RMarkdown example . . . . .	10
Brief guide to Markdown . . . . .	15
Brief guide to mathematical typesetting with L <sup>A</sup> T <sub>E</sub> X . . . . .	19
<b>Git</b>	<b>29</b>
Installation . . . . .	30
Configuration . . . . .	31
Creating and initiating a Git repository . . . . .	31
Adding and editing files . . . . .	36
Reverting changes . . . . .	39
Branching and merging . . . . .	44
Using remote repositories . . . . .	46
<b>References</b>	<b>48</b>

## Introduction

The end product of any data analysis is usually a set of tables, figures, and the seemingly countless statistics and other quantities that specify the results of the statistical modelling or testing and that was performed. These results are then usually communicated in reports, including and especially peer-reviewed scientific articles, or through talks and presentations, or through mainstream or social media, and so on. These results are the end product of an often long and arduous process, something we will call the *data analysis pipeline*, that began initially with just the raw data. For example, the original raw data might have been in the form of a set of `.xlsx` files that were downloaded from a website. These might have been wrangled and transformed repeatedly by a long series of operations, such as those described in Chapter 5, to produce various “tidy” data-sets, which themselves might then have been repeatedly and iteratively visualized and statistically analysed and modelled. It is not an exaggeration to say that this whole process might have involved hundreds or even thousands of hours of work, taking place intermittently over the course of months or even years, and involving many different people at many different stages. We can view the data analysis pipeline as akin to a factory: raw materials, in the form of the raw data, go in; these are worked on and turned into something new and valuable by the combined efforts of human labour and machines (computers); and finally end products are produced that are to be consumed by others.

The aim of *reproducible data analysis*, at its most general, is to make the data analysis factory or pipeline as open and transparent as possible, and to allow all others, including our future selves, to be able to exactly reproduce any of the results that were produced by it. In other words, the aim is to make it possible for

anyone to determine exactly where and how any given table, figure, or statistical quantity was obtained, and to be able to reproduce any of these results exactly. If the pipeline is suitably reproducible, anyone ought to be able to take the original raw data and reproduce all the final results, or alternatively ought to be able to take the raw data and analyse them in new and different ways, thus producing new and interestingly different results.

Doing reproducible data analysis is often motivated by a general commitment to doing *open science*. Since the origin of modern science in the 17th century, part of its defining ethos (see Merton 1973) has been the unrestricted sharing of the fruits of research, and also a full disclosure of all details of any research so that they may be scrutinized by others. Thus, openness and transparency are core ethical principles in all science. Recently, it has become apparent that these ethical principles, although often endorsed in principle, are not usually followed in practice, and that in fact there is a widespread culture of not sharing data (see, for example, Tenopir et al. 2011; Houtkoop et al. 2018; Fecher, Friesike, and Hebing 2015), not sharing data analysis and other computer code (see, for example, Stodden, Guo, and Ma 2013; Shamir et al. 2013), and that there is a widespread general lack of research transparency in science (see, for example, Iqbal et al. 2016). This has led to repeated calls for major cultural changes related to the sharing of data and code and general research transparency under the banner of doing *open science* (for example, Nosek et al. 2015; Munafò et al. 2017; Gorgolewski and Poldrack 2016; Ioannidis 2015).

Reproducible data analysis can also be motivated simply as a means of doing more high quality and robust data analysis, even when this analysis is not being done as part of scientific research per se, such as with confidential analyses that is done in business and industry. In these contexts, the raw data and analysis pipeline may always remain confidential and never be shared publicly. Nonetheless, doing this analysis using reproducible research practices allows for essential quality control. It allows the analysts themselves, those they do the analysis for, or future analysts who inherit the project or are brought on board it later, to scrutinize and double-check every detail of the analysis and reproduce every result. These are essential measures to identify errors, increase rigour, and verify the final results and conclusions. Even, or especially, outside of academic or scientific research, there can be enormous practical and financial incentives to minimizing errors and increasing analytical rigour in data analysis. As an example, it has been argued that a lack of reproducible data analysis techniques was partly to blame for a \$9bn loss at the investment bank JPMorgan in 2012 (see Hern 2013).

This chapter is about doing *open*, *transparent*, and *reproducible* research using R and related software tools. In particular, we will use discuss *RMarkdown* and *knitr*, *Git* for version control, using R packages as a means of sharing code and data, using build automation tools for automating computationally intensive tasks, and using *docker* for creating reproducible computing environments.

Before we proceed, however, let us briefly discuss some terminology. The terms *open*, *transparent*, and *reproducible* are a collection of adjectives that describe a set of data analysis practices. These terms are obviously not identical, but they are related. It is, however, not trivial to state the extent to which any one depends upon or requires any other. While it is not necessary to be pedantic about the definitions and scope of these terms, we will briefly outline our understanding of each term.

**Open** Open data analysis, like open source software or open science generally, is data analysis where all the data, code, and any other required materials are fully disclosed to and shared with others, usually by being made publicly available. Being publicly available, and in an unrestricted manner, is therefore usually taken to be either a defining feature of open data analysis.

**Transparent** Transparent data analysis is analysis where, as mentioned above, it is possible to determine exactly where and how any given table, figure, or statistical result was obtained. Making data and code and all other material open is usually a sufficient condition for the analysis to be transparent, but it is possible to conceive of situations where data and code are open and available but are poorly written and organized, or are obfuscated, or require undocumented manual intervention, and so lack transparency. Likewise, as also mentioned above, it is not necessary for data to be open, at least in the sense of publicly available, for it to be transparent.

**Reproducible** Reproducible data analysis is any data analysis where an independent analyst can exactly

reproduce all the results. For an analysis to be reproducible it is necessary that all the data and code is available and in full working condition. Strongly implied, however, is that running the code is essentially a turnkey operation. Software tools, such as *build automation* tools like `make`, are often used, especially with larger projects. For reports and articles, and their myriad tables and figures and in-text quantities, *literate programming* tools, particularly `rmarkdown` with R, are used. Version control software, such as `git`, is often used to organize and keep a log of the development of all the analysis code and scripts. To allow the code to be used across different operating systems and with the correct dependencies, virtual or containerized operating systems using tools such as `docker` are sometimes used. Other tools, such as the `checkpoint::checkpoint()` function in R can be used to ensure the correct versions of R package dependencies are used.

## Using RMarkdown for reproducible reports

As mentioned, the results of data analyses are communicated through reports such as peer reviewed scientific articles and other manuscripts, slide or poster presentations, webpages, and so on. For the results described in these reports to be truly reproducible, every figure, table, or any other value or quantity ought to be reproducible with minimal effort by anyone, including and especially those not directly involved in the analysis. For this, it is necessary that the data and code be made available to others. But this alone is not sufficient. Even if the necessary data and code were complete, it might still be challenging or nearly impossible for someone who was not directly involved in the analysis to identify exactly which code produced any given figure, table, or other reported quantity. This is especially the case if the results in the report were generated, as is exceedingly common, by transcribing or copying and pasting results from the output of statistical software, or inserting figures that were exported, into a document such as MS Word or  $\text{\LaTeX}$ . For any given figure or table or quantity, hunting down and verifying which piece of code produced it could be very challenging and time consuming. Moreover, there is a very high probability that the reported results will have some errors somewhere simply because they were all the product of transcription or other manual interventions. Indeed, there is no way of even categorically verifying that all the correct code and data necessary to produce the results has been made available without painstakingly finding and checking each individual piece code for each individual result.

For a report to be truly reproducible, therefore, the data, code, and resulting document need to be inextricably coupled in a manner that goes beyond how documents are traditionally written. This coupling of code, data, and text is based on the concept of *literate programming* (see Knuth 1984). In a *literate program*, the text of the documentation of a computer program and the code of that program are linked in one single file or set of files. These files can be processed by different programs either to generate the documentation manuals or the computer programs that will be compiled and executed like normal programs. While this principle has become a common means of generating documentation for computer code (see, for example, *Roxygen2* for documenting R code, *Doxygen* for documenting C++, *Sphinx* for documenting Python, etc), the same principle can be used to creating reproducible data analysis reports, which are sometimes known as *dynamic documents* (see Xie 2017) The most popular, and arguably the best, way of doing this using R is to use *RMarkdown*, or more precisely to use RMarkdown and a combination of tools including *knitr*, *pandoc*,  $\text{\LaTeX}$ , and others.

Before we proceed, let us briefly introduce or define a number of key concepts and tools.

**RMarkdown** RMarkdown itself is simply a text file format, and so an RMarkdown file is essentially a script, not unlike an R script. In fact, just like normal R scripts, we usually open, edit, and “run” RMarkdown files inside of RStudio. Unlike R scripts, which consist of just R code, possibly including comments, but nothing else, RMarkdown scripts primarily consist of a mixture of two type of code: *Markdown* code and normal R code. As we’ll see, the R code in RMarkdown documents occurs in either R *chunks*, which are simply blocks of code, or in *inline* R code, which are small amounts of R code inside Markdown code.

**Markdown** Markdown is a minimal or lightweight markup language. It was initially created by John Gruber (Gruber 2004) primarily as means of allowing web users to create formatted posts, including links,

images, lists, etc., on online discussion forums. Markdown consists of normal text, just as you would write in an email or any other document, as well some minimal syntax that instructs how this text should be formatted when it is rendered into some output document, such as a html page, pdf or Word document.

**Knitr** Knitr (see Xie 2017) is a general tool for dynamic documentation generation using R. In brief and put very simply, knitr extracts R code from RMarkdown documents, runs this code, and then, by default, inserts copies of this code and the code's output into a Markdown document. Knitr then runs *pandoc* to create a document, such as a pdf, MS Word document, etc, from this Markdown file.

**Pandoc** Pandoc (MacFarlane 2006) is, in general, a document converter that can convert a large number of input document types to an equally large number of output document types. For present purposes, it is the means by which Markdown documents generated by knitr are converted into their final output formats such as pdf, MS Word, html, etc.

**L<sup>A</sup>T<sub>E</sub>X** L<sup>A</sup>T<sub>E</sub>X [l<sup>A</sup>mp<sub>o</sub>rt1994<sub>l</sub>at<sub>e</sub>x] is a document preparation system that is specialized for creating high quality technical and scientific documents, especially those containing mathematical formulas and technical diagrams. It is widely used for creating academic and research manuscripts in mathematically oriented fields such as statistics, computer science, physics, etc. L<sup>A</sup>T<sub>E</sub>X documents are created by first writing a `.tex` source code file, which is mixture of L<sup>A</sup>T<sub>E</sub>X and T<sub>E</sub>X code, markup syntax, and plain text. This `.tex` file is then rendered to, usually, a pdf using a L<sup>A</sup>T<sub>E</sub>X rendering engine, of which there are many but the most widely is the default `pdflatex` engine.

## Installation

If using RStudio, the necessary R packages including `rmarkdown` and `knitr` are automatically installed. Likewise, the external pandoc program is also automatically installed when `rmarkdown` is installed. While this set of tools will allow you to create html and MS Word documents, L<sup>A</sup>T<sub>E</sub>X must be installed to create pdf outputs. Installation and configuration of L<sup>A</sup>T<sub>E</sub>X, because it is a large external program, may not always be straightforward. It may seem, therefore, that this L<sup>A</sup>T<sub>E</sub>X installation step is not worthwhile, especially given that many people unquestionably and happily use MS Word to write their manuscripts. However, we *highly* recommend installing L<sup>A</sup>T<sub>E</sub>X and using the L<sup>A</sup>T<sub>E</sub>X-ed pdf document output over MS Word for writing manuscripts. At the very least, the typesetting quality of the resulting L<sup>A</sup>T<sub>E</sub>X-ed document will be considerably higher than that of the MS Word word. Moreover, by using L<sup>A</sup>T<sub>E</sub>X with RMarkdown, we can avail of all the power of L<sup>A</sup>T<sub>E</sub>X to create the final document. This includes the use of all L<sup>A</sup>T<sub>E</sub>X packages for creating and styling mathematical and technical content, including technical diagrams, fine control of figures and tables, their look and feel, and their placement, internal cross referencing to refer to individual pages, sections, figures, etc., automatic index generation, and many more. Even if these features may seem unnecessary at the beginning, in our opinion, it is nonetheless still worthwhile to use L<sup>A</sup>T<sub>E</sub>X with RMarkdown from the beginning, so that all these features can be used as one's experience with RMarkdown grows.

To install L<sup>A</sup>T<sub>E</sub>X for the purposes of using with RMarkdown, we can install the R package `tinytex`.

## A minimal RMarkdown example

As just mentioned, an RMarkdown file is essentially a script containing two types of code: Markdown and normal R code. The following is the contents of small RMarkdown file, named `example.Rmd`.

```
This is some text. Some of it is *italicized*, and some is in **bold**.
```

```
```{r}
x <- c(5, 12, 103)
y <- x ^ 2
```
```

```
The mean of `x` is `r mean(x)`, and the min value of `y` is `r min(y)`.
```

This example contains one R *chunk*. This is the portion of R code between the opening ````{r}` and the closing `````. It also has two pieces of inline R code. These are the code segments between the opening ``r`` and closing ```. The remainder of the code in this example is plain Markdown code.

As we will see, we could open this file in the RStudio editor and then *knit* it to, for example, a pdf document. That pdf document would be named `example.pdf` and its contents would be, by default, as follows.

---

This is some text. Some of it is *italicized*, and some is in **bold**.

```
x <- c(5, 12, 103)
y <- x ^ 2
```

The mean of `x` is 40, and the min value of `y` is 25.

---

The key features to notice here are that the Markdown code within the RMarkdown file is formatted according to its instructions. For example, the appropriate words are italicized and emboldened. Also, the R code in the chunk is appropriately formatted as code, using a monospaced font, and is syntax highlighted. Finally, and most interestingly in this example, the *values* returned by the R commands in the two pieces of inline code are inserted at the locations of, and thus replace, the original two pieces of inline code. In other words, in the final document, the values of 40 and 25 replace ``r mean(x)`` and ``r min(y)``, respectively.

Although this is a very simple example, a lot has happened to produce the final pdf output. The procedure is roughly as follows:

- First, `knitr` extracts out the R code from `example.Rmd` and runs them in a separate R session. The code is run in the order in which it appears in the document. So, first, the R chunk is run, which creates two vectors, `x` and `y`. Then, the R commands `mean(x)` and then `min(y)` are executed.
- Next, all the Markdown code in `example.Rmd` is extracted out and inserted into a new temporary file `example.md`. Likewise, by default, copies of the R code in the chunk, though not the inline code, are inserted into this `example.md` Markdown file. Any R code that is inserted into `example.md` is marked-up as R code so that it will be properly rendered, including with syntax highlighting, in the final output document.
- Next, any *output* from the R code, whether the R chunk or the R inline code is inserted into the `example.md` file. In the example above, there is no output from the R chunk. However, obviously the commands `mean(x)` and `min(y)` both return numbers, and so these two numbers are inserted into the `example.md` Markdown file at the exact locations of where the two pieces of inline R code occurred.
- Now, `knitr` calls `pandoc`. Pandoc firsts convert the `example.md` Markdown into a L<sup>A</sup>T<sub>E</sub>X source code file named `example.tex`. This file is essentially just another script whose contents would be easily understandable to anyone who knows L<sup>A</sup>T<sub>E</sub>X.
- Next, `pandoc` calls a L<sup>A</sup>T<sub>E</sub>X rendering engine, such as `pdflatex`, and converts the L<sup>A</sup>T<sub>E</sub>X source code file `example.tex` into the pdf document `example.pdf`.
- Finally, the intermediate files, including `example.md` and `example.tex` are, by default, removed so that the only remaining files are the original RMarkdown source file `example.Rmd` and the final output document `example.pdf`.

## An extended RMarkdown example

In Figure 1, we shown the code of an RMarkdown file on the left and its corresponding pdf output on the right. This example shows many of the features of a typical RMarkdown document. We will discuss these features by looking through each section of the RMarkdown code.

### The YAML header

The first few lines of the document, specifically those lines that are delimited by the lines with the three dashes (i.e., `---`) constitute its *YAML* header.

```

---
title: "Data Analysis: A Report"
author: "Mark Andrews"
date: "October 25, 2019"
output: pdf_document
---

```{r setup, echo=FALSE}
knitr::opts_chunk$set(message = F,
                        warning = F,
                        out.width = "45%",
                        fig.align='center')
...

# Introduction

First, we will load the `tidyverse` packages, and read
in the the from a `.csv` file.
```{r load_packages_data}
library(tidyverse)
data_df <- read_csv('example.csv')
...

# Analysis

Here, we do a Pearson's correlation analysis.
```{r analysis}
(corr_model <- cor.test(~ x + y, data = data_df))
...

The correlation coefficient is
`r round(corr_model$estimate, 3)` .

# Visualization

The scatterplot between $x$ and $y$ is shown
in Figure \ref{fig:vis}.

```{r vis, echo=F, fig.cap='A scatterplot.'}
ggplot(data_df, aes(x, y)) +
  geom_point() +
  theme_classic()
...

```

## Data Analysis: A Report

Mark Andrews

October 25, 2019

### Introduction

First, we will load the `tidyverse` packages, and read in the the from a `.csv` file.

```
library(tidyverse)
data_df <- read_csv('example.csv')
```

### Analysis

Here, we do a Pearson's correlation analysis.

```
(corr_model <- cor.test(~ x + y, data = data_df))
```

```
##
## Pearson's product-moment correlation
##
## data: x and y
## t = 2.1583, df = 48, p-value = 0.03593
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
##  0.02080288 0.53175304
## sample estimates:
##      cor
## 0.2974284
```

The correlation coefficient is 0.297.

### Visualization

The scatterplot between  $x$  and  $y$  is shown in Figure 1.

1

Figure 1: An example of a RMarkdown file on the left, and its corresponding pdf output on the right. The pdf output is generated by *knitting* the RMarkdown file. In RStudio, this can be accomplished by `rmarkdown::render("example-2.Rmd")`, where `example-2.Rmd` is the RMarkdown file.

```

---
title: "Data Analysis: A Report"
author: "Mark Andrews"
date: "October 25, 2019"
output: pdf_document
---

```

YAML (a recursive acronym for *Y*AML *A*in't a *M*arkup *L*anguage) is itself a minimal markup language that is now often used for software configuration files. As is perhaps clear from this example, in this header we specify the title, author, date, and output format of the document. This information is used when rendering the output document. At its simplest, YAML consists of *key-value* mappings where the term on the left of the `:` is the key, and the term or statement on the right is the value. Thus, `author: "Mark Andrews"` indicates that the `author` is `"Mark Andrews"`, and when the final document is being generated, the `"Mark Andrews"` will be inserted as the value of the `author` in output document's template. Note that although we use double quotation marks as the values of the `title`, `author`, and `date` keys, neither single nor double quotation marks are always necessary. In this example, we would only require the quotations for the value of `title`. This is to ensure that the colon within `Data Analysis: A Report` is not mistakenly parsed by the YAML parser as indicating a key-value mapping.

We specify the document output type in the header by `output: pdf_document`. As mentioned above, when writing manuscripts, we recommend always using the pdf based output document option as this will be generated by L<sup>A</sup>T<sub>E</sub>X. However, had we preferred to create an MS Word output document, we would state `output: word_document`. Likewise, had we preferred a HTML output document, which can obviously be viewed in any browser or used to create a webpage, we would do `output: html_document`.

## The setup R chunk

It is common, though not necessary to have a setup R chunk like the following at the beginning of our RMarkdown file.

```
```{r setup, echo=FALSE}
knitr::opts_chunk$set(message = F,
                        warning = F,
                        out.width = "45%",
                        fig.align='center')
```
```

In addition to including a common configuration statement, this example also has some important general RMarkdown features. First, in the minimal RMarkdown example in the previous statement, the R chunk was defined by starting with the opening ````{r}` and ending with the closing `````. Here, the chunk statement begins with

```
```{r setup, echo=FALSE}
```

The term `setup` after the initial `r` is simply a label for this chunk. It is not necessary to have a label for chunks but it can be used for file navigation purposes in the RStudio editor, and is used for cross referencing purposes, especially for so-called *floating* figures and tables, as we will see. As such, it is probably a good habit to always use a (unique) label for each R chunk.

After the chunk label, we have the chunk configuration option `echo=FALSE`. This indicates the the R code in this chunk should not be displayed in the output document. Next, we have the following R statement.

```
knitr::opts_chunk$set(message = F,
                        warning = F,
                        out.width = "45%",
                        fig.align='center')
```

Here, we are set the values of `opts_chunk` list of global options used by `knitr`. For example, in this case, we start by indicating that, by default, all subsequent chunks in the RMarkdown document should have the options: `message = FALSE` and `warning = FALSE` (because these are R statements, we can use `F` for `FALSE`). Setting `message = FALSE` entails that the messages, which can be verbose, produced by R commands are not shown. As an example of this, consider the usual output when we load a package, in this case using `lme4`.

```
library(lme4)
```

```
## Loading required package: Matrix
##
## Attaching package: 'Matrix'
##
## The following objects are masked from 'package:tidyr':
##
##     expand, pack, unpack
```

Therefore, globally setting `message = FALSE` can considerably reduce clutter in the output document. Likewise, it is often preferable to globally suppress R warnings to avoid clutter in the final document. Although not used in this example, another commonly set global option at this point is `echo=FALSE`. Often, especially in manuscripts for peer-reviewed articles, we do not wish to display all, or even any, of the R code

used in the analysis. Rather, we just want to display the results of the analysis in tables, figures, etc, and so globally setting `echo=FALSE` avoids having to set `echo=FALSE` on each subsequent chunk.

The next two settings, `out.width = "45%"` and `fig.align='center'`, pertain to how figures should be displayed by default. The `out.width = "45%"` indicates that the figure should occupy 45% of the width of the page, and `fig.align='center'` indicates that it should be centered on the page.

## Markdown sections

The use of the `#` followed by some text at the start of a line in the Markdown language indicates a section header. Thus, in our case, the following line indicates we should have a section entitled *Introduction* in our document.

```
# Introduction
```

While the single `#` is for a section, subsections are indicated by `##`, and subsubsections are indicated by `###`. Here is an example with three levels of sections.

```
# This is a section
```

```
## This is a subsection
```

```
# This is another section
```

```
## This is a subsection
```

```
### This is a subsubsection
```

In principle, further nested subsections, e.g. subsubsubsections and so on, are possible. However, this does depend on both the output document type and also the templates for this documents that are used by `knitr` or `pandoc`.

## R chunk outputs

Just as we see the output of R commands when those commands are run in the console, the output of R commands will appear in the output document unless we specify request otherwise. Consider the following lines from our example RMarkdown file.

```
```${r analysis}
(corr_model <- cor.test(~ x + y, data = data_df))
```
```

In this case, the assignment statement is surrounded by `(` and `)` and in general in R, this causes the output of the expression in this statement to be shown. In other words, the output that will be shown will be identical to what would be shown had be just typed.

```
cor.test(~ x + y, data = data_df)
```

By default, this output will have each line beginning with `##` (because they are part of the R chunk's output are not interpreted as subsection headers). Should we wish use alternative symbols at the beginning of the output, we can indicate this by setting the value of the `comment` chunk option. For example, if we wanted all R output to start with `'>'`, we set `comment = '>'`, as in the following code.

```
```${r, echo=T, comment='>'}
rnorm(5) %>% round(2)
```
```

This code would then be rendered as follows.



```
rnorm(5) %>% round(2)
```

```
> [1] -0.26  0.12  0.83 -0.17  0.24
```

## L<sup>A</sup>T<sub>E</sub>X mathematical typesetting

On the following line

The scatterplot between  $x$  and  $y$  is shown

we have the terms  $x$  and  $y$ . The dollar symbols indicate that  $x$  and  $y$  should be parsed using L<sup>A</sup>T<sub>E</sub>X and typeset according. In this example, this will simply make the  $x$  and  $y$  be shown in a different and italicized font compared to normal letters. However, in general, there is a vast number of mathematical notations, formulae, and symbols that can be used here. For example, consider the following Markdown code.

If  $\Phi = \phi_1, \phi_2 \dots \phi_k \dots \phi_k$ , where each  $0 \leq \phi_k \leq 1$ , and  $\sum_{k=1}^K \phi_k = 1$ , then  $\Phi$  is a probability mass function.

This code would then be rendered as follows.

If  $\Phi = \phi_1, \phi_2 \dots \phi_k \dots \phi_k$ , where each  $0 \leq \phi_k \leq 1$ , and  $\sum_{k=1}^K \phi_k = 1$ , then  $\Phi$  is a probability mass function.

In these examples, all the L<sup>A</sup>T<sub>E</sub>X code occurs in its *inline mode*, and this obtained by using  $\$$  delimiters. In addition to inline mode, there is L<sup>A</sup>T<sub>E</sub>X *\*display\** mode. This is obtained by using  $\$$  delimiters. L<sup>A</sup>T<sub>E</sub>X display mode is used to display mathematical formulae and other notation on newlines. Consider the following example.

The probability of the observed data is as follows:

```
$$
\mathrm{P}(x_1 \dots x_n \text{vert } \mu, \sigma^2)
= \prod_{i=1}^n \mathrm{P}(x_i \text{vert } \mu, \sigma^2),
$$
where  $\mu$  and  $\sigma$  are parameters.
```

This would then be rendered as follows.

---

The probability of the observed data is as follows:

$$P(x_1 \dots x_n | \mu, \sigma^2) = \prod_{i=1}^n P(x_i | \mu, \sigma^2),$$

where  $\mu$  and  $\sigma$  are parameters.

---

L<sup>A</sup>T<sub>E</sub>X provides a very large number of mathematical symbols, notations, and formulae. It is beyond the scope of this chapter to even provide an overview of all of these features. However, we will provide an overview of some of the more widely used examples in a later section of this chapter.

## Figures

If the R code in the a chunk generates a figures, for example by using `ggplot`, then that figure is inserted into the document immediately after the location of the chunk. These figures can optionally have figure captions by setting the value of `fig.cap` in the chunk header. When the `fig.cap` is set, at least when using the `pdf_document` output, the figures then *floats*. A floating figure is one whose placement location in the final document is determined by an algorithm that attempts to minimize a large number of constraints, such as the amount of whitespace on the page, the number of figures on one page, and so on.

In our example file, we have the following lines.

```

```{r vis, echo=F, fig.cap='A scatterplot.'}
ggplot(data_df, aes(x, y)) +
  geom_point() +
  theme_classic()
```

```

This chunk will create a `ggplot` figure that will float. Because of the global settings of `out.width = 40%` and `fig.align = 'center'`, the figure will be relatively small and centered on the page. Because there is sufficient space, the figure will be placed at the bottom of the page. However, were we to increase the width of the figure to even `out.width = 50%`, the default settings of the float placement algorithm would place the figure alone on the following page. Often, the choices made by L<sup>A</sup>T<sub>E</sub>X's float placement algorithm are initially unappealing. The recommendation, however, is to not attempt any fine control until the document is complete. As each new text block, or figure, or table is added to the document, the floats are relocated. When the document is complete, it is possible to influence float placement by changing parameter settings on the float placement algorithm.

It should be noted that if figures do not float, they are placed exactly where the chunk's output would otherwise occur. In some cases, especially if the figures are relatively small, this can be satisfactory.

## An additional RMarkdown example

```

---
title: "Data Analysis: Report II"
author: "Mark Andrews"
date: "October 25, 2019"
output:
  pdf_document:
    keep_tex: yes
header-includes:
- \usepackage{booktabs}
bibliography: refs.bib
biblio-style: apalike
---

```{r setup, echo=FALSE}
knitr::opts_chunk$set(warning = FALSE, message = FALSE)
```

```{r load_packages, echo = FALSE}
library(tidyverse)
library(knitr)
library(kableExtra)

weight_df <- read_csv('weight.csv')
weight_df_grouped <- group_by(weight_df, gender)

sample_size <- weight_df_grouped %>%
  summarise(n=n()) %>%
  deframe()

weight_df_summary <- weight_df_grouped %>%
  summarise_at(vars(weight, height),
    list(avg=mean, stdev=sd))
```

# Descriptive statistics

In this data set, we have measured the weight (in kg)
and height (in cm) of `r sum(sample_size)` participants
(`r sample_size['Male']` males, `r sample_size['Female']`
females). In the following table, we show the mean and

separately for males and females.

```{r descriptives_table, echo=F}
weight_df_summary %>%
  kable(format = "latex",
    booktabs = TRUE,
    digits = 2,
    align = 'c') %>%
  kable_styling(position = "center")
```

# Statistical model

We will model the relationship between weight and height
as a varying intercepts normal linear model as follows.
For each  $i$  in 1 \dots  $n$ ,

$$y_i \sim N(\mu_i, \sigma^2), \backslash$$


$$\mu_i \sim \beta_0 + \beta_1 x_i + \beta_2 z_i,$$

where  $y_i$ ,  $x_i$ ,  $z_i$  are the weight, height,
and gender of participant  $i$ .

In R, this analysis can be easily performed as follows.
```{r stat_model}
model <- lm(weight ~ height + gender, data = weight_df)
```
```{r model_results, echo=F}
R_sq <- summary(model)$r_sq
f_stat <- summary(model)$fstatistic
p_value <- pf(f_stat[1], f_stat[2], f_stat[3], lower.tail = F)
```

The  $R^2$  for this model is `r round(R_sq, 2)`,
 $F(\text{r round}(f\_stat[2]), \text{r round}(f\_stat[3])) =$ 
`r round(f_stat[1], 2)`,
 $p$  `r format.pval(p_value, eps = 0.01)`$.

More information about varying intecept models can be
found in @GelmanHill:2007.

# References

```

Figure 2: An additional RMarkdown file example. The rendered pdf output document is shown in Figure 3.

In Figure 2, we provide the code of another RMarkdown file. This example provides some additional important

# Data Analysis: Report II

Mark Andrews

October 25, 2019

## Descriptive statistics

In this data set, we have measured the weight (in kg) and height (in cm) of 6068 participants (4082 males, 1986 females). In the following table, we show the mean and the standard deviations of the weights and heights, separately for males and females.

| gender | weight_avg | height_avg | weight_stdev | height_stdev |
|--------|------------|------------|--------------|--------------|
| Female | 67.76      | 162.85     | 10.98        | 6.42         |
| Male   | 85.52      | 175.62     | 14.22        | 6.86         |

## Statistical model

We will model the relationship between weight and height as a varying intercepts normal linear model as follows. For each  $i \in 1 \dots n$ ,

$$y_i \sim N(\mu_i, \sigma^2),$$
$$\mu_i = \beta_0 + \beta_1 x_i + \beta_2 z_i,$$

where  $y_i$ ,  $x_i$ ,  $z_i$  are the weight, height, and gender of participant  $i$ .

In R, this analysis can be easily performed as follows.

```
model <- lm(weight ~ height + gender, data = weight_df)
```

The  $R^2$  for this model is 0.45,  $F(2, 6065) = 2495.83$ ,  $p < 0.01$ .

More information about varying intercept models can be found in Gelman and Hill (2007).

References

Figure 3: The pdf document output corresponding to the RMarkdown code shown in Figure 2.

RMarkdown features that were not covered in the previous example. The pdf output document corresponding to this document is shown in Figure 3.

## Additional YAML header options

In this example file, we have set some additional settings in the YAML header. For example, we have the following lines.

```
output:
  pdf_document:
    keep_tex: yes
header-includes:
- \usepackage{booktabs}
```

The latter two lines relate to the bibliography, and we will return to these below. The first two lines indicates that the  $\LaTeX$  command `\usepackage{booktabs}` should be included in the header of the resulting `.tex` file prior to be rendered to pdf by the  $\LaTeX$  engine. Recall that as mentioned above, the rendering process to produce the final `.pdf` document is that `knitr` takes the `.Rmd` file and creates a `.md` file, with R code and R code output inserted into it. If the final output is to be pdf, `pandoc` converts the `.md` file to a `.tex` file, and that is then rendered to `.pdf` by a  $\LaTeX$  engine such as `pdflatex`. When `pandoc` creates the `.tex` file from the `.md`, it uses `.tex` file templates that load some of the more widely used  $\LaTeX$  packages. However, if any other  $\LaTeX$  packages are required to create the final pdf document, they may be listed as we have done here in the YAML header using the `header-includes` option. It should be noted that there are many thousands of additional  $\LaTeX$  packages that can be installed.

Another option to include additional  $\text{\LaTeX}$  functionality in your RMarkdown is to include a block of  $\text{\LaTeX}$  code into the header. For example, we might have a file named `include.tex` with the following  $\text{\LaTeX}$  command.

```
\newcommand{\Prob}[1]{\mathrm{P}( #1 )}
```

If we had the following YAML header in our RMarkdown document, the commands in `include.tex` would be available to use in the body of the RMarkdown file.

```
output:
  pdf_document:
    includes:
      in_header: include.tex
```

For example, in this case, we could now write `\Prob{X = x} = 0.25` in our RMarkdown file, and this would eventually be rendered as  $P(X = x) = 0.25$ . The file `include.tex` can have an extensive number of custom commands, including numerous `\usepackage{}` statements.

## Formatted tables

When writing reports, we often wish to present results of statistical analyses in tables. Naturally, we would ideally like these tables to be formatted to a high standard. As we've seen, we can always simply display the standard output of R commands that produce tables and data frames, but the resulting unformatted monospaced font text is not of an acceptable standard for most manuscripts and reports, especially those that are intended for eventual publication. As an example, consider the table produced by the following R commands that use the built in `swiss` data set.

```
(swiss_df <- swiss %>%
  select(Fertility:Examination) %>%
  slice(1:5))
```

```
##      Fertility Agriculture Examination
## 1         80.2          17.0           15
## 2         83.1          45.1            6
## 3         92.5          39.7            5
## 4         85.8          36.5           12
## 5         76.9          43.5           17
```

Clearly, this output is not sufficient for all but informal reports.

Fortunately, we have a number of options for formatting tables. For example, to use `kable` to create a table specifically for use in the `pdf_document` output, we can do the following.

```
library(knitr)
swiss_df %>%
  kable(format = 'latex',
        booktabs = TRUE,
        align = 'c')
```

| Fertility | Agriculture | Examination |
|-----------|-------------|-------------|
| 80.2      | 17.0        | 15          |
| 83.1      | 45.1        | 6           |
| 92.5      | 39.7        | 5           |
| 85.8      | 36.5        | 12          |
| 76.9      | 43.5        | 17          |

Here, we use `booktabs = TRUE` and in the YAML header, as mentioned, we ensured that the `booktabs`  $\text{\LaTeX}$  package is loaded. By loading the `booktabs` package, the typesetting of the tables in  $\text{\LaTeX}$  is often improved over the default options. The `align = 'c'` option ensures that the values in each column in the

table are centered. If we want to center the table produced by the above `kable` command, we can use the `kable_styling` function in the `kableExtra` package as follows.

```
library(kableExtra)
swiss_df %>%
  kable(format = 'latex',
        booktabs = TRUE,
        align = 'c') %>%
  kable_styling(position = 'center')
```

| Fertility | Agriculture | Examination |
|-----------|-------------|-------------|
| 80.2      | 17.0        | 15          |
| 83.1      | 45.1        | 6           |
| 92.5      | 39.7        | 5           |
| 85.8      | 36.5        | 12          |
| 76.9      | 43.5        | 17          |

## Display maths

In this example, we provide some examples of *display* mode mathematics. As mentioned, this is where the mathematical statements are on separate lines. In this example, we use the `aligned` environment in  $\text{\LaTeX}$ , which allows us to align multiple mathematical statements. Specifically, the example we use is as follows.

```
$$
\begin{aligned}
y_i &\sim N(\mu_i, \sigma^2), \\
\mu_i &= \beta_0 + \beta_1 x_i + \beta_2 z_i,
\end{aligned}
$$
```

This is rendered as follows.

$$y_i \sim N(\mu_i, \sigma^2),$$

$$\mu_i = \beta_0 + \beta_1 x_i + \beta_2 z_i,$$

The `&` symbol on each line is used to align the lines, so that the  $\sim$  on the first line is aligned with the  $=$  on the second line. Here, we need to add `\\` at the end of the first line to force a line break. The `aligned` environment is widely used in  $\text{\LaTeX}$ , especially for showing algebraic derivations. For example, the following

```
$$
\begin{aligned}
y &= \frac{e^x}{1+e^x}, \\
&= \frac{e^x}{e^x(e^{-x} + 1)}, \\
&= \frac{1}{1+e^{-x}}
\end{aligned}
$$
```

is rendered as follows.

$$y = \frac{e^x}{1+e^x},$$

$$= \frac{e^x}{e^x(e^{-x} + 1)},$$

$$= \frac{1}{1+e^{-x}}.$$

## Formatted inline R output

In this RMarkdown example, we display the  $R^2$  of the linear model as well as its accompanying F statistic and p-value with the following code.

```
The  $R^2$  for this model is `r round(R_sq, 2)`,  
 $F(\text{`r round(f_stat[2])`}, \text{`r round(f_stat[3])`}) =$   
`r round(f_stat[1], 2)` $,  
 $p$  `r format.pval(p_value, eps = 0.01)` $.
```

Here, we make use of some quantities, i.e., `R_sq`, `f_stat`, `p_value`, created earlier in the file. In addition, we place some of inline R code inside `$` so that they will be formatted as mathematical notation. Placing inline R code inside either  $\text{\LaTeX}$  mathematical environments is always possible. Consider the following example.

```
$$  
\int_{-\infty}^1 \frac{1}{\sqrt{2\pi}} e^{-\frac{x^2}{2}} dx = `r round(pnorm(1), 3)`  
$$
```

This will produce the following output.

$$\int_{-\infty}^1 \frac{1}{\sqrt{2\pi}} e^{-\frac{x^2}{2}} dx = 0.841$$

Note that also in this example, we also avail of the R function `format.pval`. This function can be used to return an inequality when the p-value is lower than a certain threshold named `eps`. For example, consider the following example.

```
p <- c(0.05, 0.02, 0.011, 0.005, 0.001)  
format.pval(p, eps = 0.01)
```

```
## [1] "0.050" "0.020" "0.011" "<0.01" "<0.01"
```

## Bibliography

In the YAML header, we have included the following statements.

```
bibliography: refs.bib  
biblio-style: apalike
```

The first line tells RMarkdown to use the bibliographic information contained in the file `refs.bib`, which is assumed to be present in the working directory. The second line tells RMarkdown to format the citations and the references in the bibliography using APA style.

The `refs.bib` file is just a plain text file, and can be named anything. In this example, its content is minimal and contains just the following content.

```
@book{GelmanHill:2007,  
  address = {New York},  
  author = {Gelman, Andrew and Hill, Jennifer},  
  title = {Data Analysis Using Regression and Multilevel/Hierarchical Models},  
  publisher = {Cambridge University Press},  
  year = 2007  
}
```

The `@book{ .... }` content is a  $\text{BIB}\text{\TeX}$  bibliographic entry for a book, with  $\text{BIB}\text{\TeX}$  being the primary bibliography manager used with  $\text{\LaTeX}$ . As this is a simple text file, these bibliographic entries can be relatively easily created. However, because  $\text{BIB}\text{\TeX}$  is such a widely used bibliographic manager,  $\text{BIB}\text{\TeX}$  bibliographic entries are provided in many scholarly article data bases, including Google Scholar. It is therefore very easy to build up a personal bibliography file, which can then be re-used in all of one's reports.

Having defined a bibliography in the YAML header as we have done, we can now use the `BIBTEX` *keys* to perform citations. The *key* for `BIBTEX` entry is the text after the first `{` and before the first comma. Consider the following 3 `BIBTEX` entries.

```
@book{xie2017dynamic,
  title={Dynamic Documents with R and knitr},
  author={Xie, Yihui},
  year={2017},
  publisher={Chapman and Hall/CRC}
}
```

```
@article{knuth1984literate,
  title={Literate programming},
  author={Knuth, Donald Ervin},
  journal={The Computer Journal},
  volume={27},
  number={2},
  pages={97--111},
  year={1984},
  publisher={Oxford University Press}
}
```

```
@book{wickham:2019,
  Author = {Hadley Wickham},
  Title = {Advanced R},
  Publisher = {Chapman and Hall/CRC},
  Year = {2019},
  edition = 2
}
```

The keys here are `xie2017dynamic`, `knuth1984literate`, and `wickham:2019`. Should we wish to refer to, say, Wickham’s book in our RMarkdown file, we would simply write `@wickham:2019`. For example, the following statement in RMarkdown

```
As is described in @wickham:2019, R is a functional programming language.
```

would result in “As in described in Wickham (2019), R is a functional programming language.” in the output document. Alternatively, had we written

```
R is a functional programming language [ @wickham:2019 ].
```

this would result in “R is a functional programming language (Wickham 2019).” in the output document.

In either case, the following line would then be inserted at the end of the output document.

Wickham, Hadley. 2019. *Advanced R*. 2nd ed. Chapman; Hall/CRC

If we had multiple citations, they would all be listed at the end of the output document in alphabetical order. By default in RMarkdown, this list of bibliographic references are not given a section name. Therefore, as we have done in this example RMarkdown document, we simply end the document with a section named *References* by using the following line as the last line.

```
# References
```

All the references will now be inserted after this section header. Should we prefer the name *Bibliography*, or another other option here, we simply change the name of this section.

## Brief guide to Markdown

As a minimal markup language, there are not too many

## Headers

```
# Introduction

This is a sentence in the Introduction.

## Objectives of study

This is a sentence in the first subsection
of the Introduction.

# Analysis

This is a sentence in the Analysis section.

## Exploratory analysis

This is a sentence of the first subsection
of the Analysis section.

## Statistical model

This is a sentence of the second subsection
of the Analysis.

# Conclusion

This is a sentence in the last section
of the document.
```

## Introduction

This is a sentence in the Introduction.

### Objectives of study

This is a sentence in the first subsection of the Introduction.

## Analysis

This is a sentence in the Analysis section.

### Exploratory analysis

This is a sentence of the first subsection of the Analysis section.

### Statistical model

This is a sentence of the second subsection of the Analysis.

## Conclusion

This is a sentence in the last section of the document.

Figure 4: An example of a Markdown file on the left, and its corresponding pdf output on the right. The Markdown file has sections, indicated by lines beginning with #, and subsections, indicated by lines beginning with ##.

We have already seen section headers above. Section headers begin with the # at the start of a new line followed by the section title. These lines are preceded and followed by a blank line. The number of # symbols indicate the level of the section: # indicates a section, ## is a subsection, ### is a subsubsection, and so on. Note that there should be a space after the # symbol or symbols at the start of the line. For example, # Introduction will create a section entitled *Introduction*, but #Introduction will produce a line with #Introduction on it. In Figure 4, on the left, we see a Markdown file with multiple sections and subsections. On the right, we see this file rendered as a pdf document.

## Font style and weight

As we've seen above in a few examples, we can produce italicized text by surrounding the text with \*, and we obtain bold text by surrounding it with \*\*. Here are some examples.

- `*this is italicized text*` produces *this is italicized text*
- `**this is bold text**` produces **this is bold text**
- `***this is bold and italicized text***` produces ***this is bold and italicized text***

Underscores, i.e. \_, have the same effect as \*.

- `_this is italicized_` produces *this is italicized*
- `__this is in bold__` produces **this is in bold**
- `___this is in bold and in italics___` produces ***this is in bold and in italics***

We can also mix \_ and \*.

- `*_this is italicized bold text_*` produces ***this is italicized bold text***
- `__*this is italicized bold text*_` produces ***this is italicized bold text***
- `_**this is italicized bold text**_` produces ***this is italicized bold text***



In Markdown, there is no general way of producing underlined text. However, if we are using `pdf_document` output, we can use the  $\text{\LaTeX}$  `\underline` command. For example, `\underline{this text is underlined}` gives this text is underlined.

## Code

Often in technical documents, we need to display computer code. If we simply surround the code block by `````, we will obtain monospace typed text. For example, the following Markdown code shows Python code.

```
```  
for x_i in x:  
    y.append(x_i)  
```
```

This is rendered as follows.

```
for x_i in x:  
    y.append(x_i)
```

However, ideally we would prefer the code to be syntax highlighted. We can do so by indicating the code's language after the initial `````. In the following Markdown code, we state that the code is Python.

```
```python  
for x_i in x:  
    y.append(x_i)  
```
```

This is then rendered as follows.

```
for x_i in x:  
    y.append(x_i)
```

When we wish to display R code, we have another option. We can use a normal R code chunk, but set the chunk parameter `eval` to `FALSE`, and set `echo` to `TRUE`, assuming it is not globally set to `TRUE`.

```
```{r, echo=TRUE, eval=FALSE}  
n <- 10  
x <- rnorm(n)  
y <- 2.25 * x + rnorm(n)  
  
M <- lm(y ~ x)  
```
```

This is then rendered as follows.

```
n <- 10  
x <- rnorm(n)  
y <- 2.25 * x + rnorm(n)  
  
M <- lm(y ~ x)
```

## Links and cross references

### Lists

There are three types of lists that are possible with the pdf document output: Itemized lists, enumerated lists, and definition lists.

**Itemized lists** In itemized lists, also known as unordered lists or even bullet-point lists, each item begins on a new line that begins with `*` followed by a space. The list has to be both preceded and is followed by a blank line. For example, we have Markdown code with an listed list on the left, and its rendered output on the right.

```
* Apple
* Orange
* Blueberry
```

- Apple
- Orange
- Blueberry

Items in a list do not need to be written using only one line of Markdown code. A single item may be spread over multiple lines as in the following example.

```
* If we start our item
on one line,
and continue to another
line with no break,
it will still appear as one item.
* The same thing happens
  if you continue
    on the next line after
    and indention.
```

- If we start our item on one line, and continue to another line with no break, it will still appear as one item.
- The same thing happens if you continue on the next line after and indention.

We can also create nested lists by using indented lines that themselves begin with `*`, as in the following example.

```
* Farm animal
  * Cow
  * Pig
  * Sheep
* Wild animal
  * Fox
  * Wolf
```

- Farm animal
  - Cow
  - Pig
  - Sheep
- Wild animal
  - Fox
  - Wolf

In general in Markdown, the indentation should be made with four spaces or a tab. In some cases, fewer spaces than 4 will suffice, but it is generally better to use four spaces. Also, beware that some editors, including the RStudio editor, map the Tab key to 2 spaces.

**Enumerated lists** Enumerated lists are defined and behave just like itemized lists except that instead of an item or sub-item defined by a line beginning with a `*` (followed by a space), it is defined by a line beginning a number followed by a `.` and then space. In the following example, our items begin with `1.` followed by a space.

```
1. Potato
1. Broccoli
1. Cabbage
```

1. Potato
2. Broccoli
3. Cabbage

We do not need to always use `1.` to obtain an enumerated list. Any other numbers will suffice, as in the following example.

```
1. Potato
4. Broccoli
10. Cabbage
```

1. Potato
2. Broccoli
3. Cabbage

However, if we use a number greater than 1 as the first item, then that will be used as the starting value of the numbering of the list as in the following example.

```
12. Potato
4. Broccoli
10. Cabbage
```

12. Potato
13. Broccoli
14. Cabbage

Enumerated lists can be nested just like itemized lists.

1. Tree
  1. Fir
  1. Pine
  1. Oak
1. Flower
  1. Rose
  1. Tulip
  1. Daffodil

1. Tree
  1. Fir
  2. Pine
  3. Oak
2. Flower
  1. Rose
  2. Tulip
  3. Daffodil

We may also mix enumerated and itemized lists when using nestings as in the following example.

- \* Tree
  1. Fir
  1. Pine
  1. Oak
- \* Flower
  1. Rose
  1. Tulip
  1. Daffodil

- Tree
  1. Fir
  2. Pine
  3. Oak
- Flower
  1. Rose
  2. Tulip
  3. Daffodil

**Definition lists** Definition lists, when rendered, begin with an emboldened term followed by a definition or description. This can be useful for definitions per se, but also when we need to elaborate on or describe certain terms. For example, they could be used to describe the meaning of different variables in a data set. We create a definition list by beginning a new line with some text, which is usually brief, such as a name. The subsequent line then begins with a `:` followed by the definition or description. For example, in the following code, we provide a definition list elaborating upon the variables in a dataset.

```
participant-id
: An integer that uniquely
  codes each participant.

gender
: A binary variable with
  values `female` and `male`.

age
: A numeric variable giving
  the participant's age in years.
```

```
participant-id An integer that uniquely codes
each participant.

gender A binary variable with values female and
male.

age A numeric variable giving the participant's
age in years.
```

## Brief guide to mathematical typesetting with L<sup>A</sup>T<sub>E</sub>X

RMarkdown provides an extensive set of commands for typesetting mathematical formulas, symbols, and other notation. These commands are in fact L<sup>A</sup>T<sub>E</sub>X commands, but are available regardless of the output document format. In other words, even if we are not using the `pdf_document` output format, which is ultimately rendered by a L<sup>A</sup>T<sub>E</sub>X typesetting engine, we may still avail of these L<sup>A</sup>T<sub>E</sub>X commands for mathematical typesetting.

As we have seen above, there are two *modes* in which mathematical formulas and notation can appear: *inline* mode, and *display* mode. Inline model is where the mathematical notation appears within a line of normal text. It is obtained by surrounding the code a `$`. For example, the following code

```
Einstein's famous formula for mass-energy equivalence is $E = mc^2$.
```

is rendered as follows:

Einstein's famous formula for mass-energy equivalence is  $E = mc^2$ .

On the other hand, display mode is where the mathematical formulas or notation appear on a line on their own. It is obtained by surrounding the code by `$$`. For example, the following code

```
Einstein's famous formula for mass-energy equivalence is
$$
E = mc^2.
$$
```

is rendered as follows:

Einstein's famous formula for mass-energy equivalence is

$$E = mc^2.$$

In what follows, we provide a guide to the main mathematical typesetting commands that are available in RMarkdown. However, we aim to provide no more than a brief or minimal introduction. To able to appreciate all of the mathematical typesetting options in RMarkdown, we would have to provide a thorough introduction to mathematical typesetting in L<sup>A</sup>T<sub>E</sub>X. This is well beyond the scope of this book, but there are many books that provide extensive details about mathematical typesetting in L<sup>A</sup>T<sub>E</sub>X. For example, one highly recommended comprehensive and up to date book is Grätzer (2016).

## Symbols for variables

In mathematical formulas and notation, there are a large number of symbols that are commonly used as variables. These particularly include the upper and lower case letters of the English alphabet, most upper and lower case letters of the Greek alphabet, and some Hebrew letters. In addition, some special fonts, namely calligraphic and *blackboard bold* fonts, of English uppercase letters commonly appear too.

The English upper and lowercase letters appear in italics font in either inline or display mode, as we see in the following table.

|   |          |   |          |   |          |   |          |   |          |   |          |   |          |
|---|----------|---|----------|---|----------|---|----------|---|----------|---|----------|---|----------|
| a | <i>a</i> | i | <i>i</i> | q | <i>q</i> | y | <i>y</i> | G | <i>G</i> | O | <i>O</i> | W | <i>W</i> |
| b | <i>b</i> | j | <i>j</i> | r | <i>r</i> | z | <i>z</i> | H | <i>H</i> | P | <i>P</i> | X | <i>X</i> |
| c | <i>c</i> | k | <i>k</i> | s | <i>s</i> | A | <i>A</i> | I | <i>I</i> | Q | <i>Q</i> | Y | <i>Y</i> |
| d | <i>d</i> | l | <i>l</i> | t | <i>t</i> | B | <i>B</i> | J | <i>J</i> | R | <i>R</i> | Z | <i>Z</i> |
| e | <i>e</i> | m | <i>m</i> | u | <i>u</i> | C | <i>C</i> | K | <i>K</i> | S | <i>S</i> |   |          |
| f | <i>f</i> | n | <i>n</i> | v | <i>v</i> | D | <i>D</i> | L | <i>L</i> | T | <i>T</i> |   |          |
| g | <i>g</i> | o | <i>o</i> | w | <i>w</i> | E | <i>E</i> | M | <i>M</i> | U | <i>U</i> |   |          |
| h | <i>h</i> | p | <i>p</i> | x | <i>x</i> | F | <i>F</i> | N | <i>N</i> | V | <i>V</i> |   |          |

If we prefer to have non-italicized, i.e. upright or Roman, fonts in math maths, we can use the command `\mathrm` as we see in the following table.

|                         |   |                         |   |                         |   |                         |   |                         |   |
|-------------------------|---|-------------------------|---|-------------------------|---|-------------------------|---|-------------------------|---|
| <code>\mathrm{a}</code> | a | <code>\mathrm{m}</code> | m | <code>\mathrm{y}</code> | y | <code>\mathrm{K}</code> | K | <code>\mathrm{W}</code> | W |
| <code>\mathrm{b}</code> | b | <code>\mathrm{n}</code> | n | <code>\mathrm{z}</code> | z | <code>\mathrm{L}</code> | L | <code>\mathrm{X}</code> | X |
| <code>\mathrm{c}</code> | c | <code>\mathrm{o}</code> | o | <code>\mathrm{A}</code> | A | <code>\mathrm{M}</code> | M | <code>\mathrm{Y}</code> | Y |
| <code>\mathrm{d}</code> | d | <code>\mathrm{p}</code> | p | <code>\mathrm{B}</code> | B | <code>\mathrm{N}</code> | N | <code>\mathrm{Z}</code> | Z |
| <code>\mathrm{e}</code> | e | <code>\mathrm{q}</code> | q | <code>\mathrm{C}</code> | C | <code>\mathrm{O}</code> | O |                         |   |
| <code>\mathrm{f}</code> | f | <code>\mathrm{r}</code> | r | <code>\mathrm{D}</code> | D | <code>\mathrm{P}</code> | P |                         |   |
| <code>\mathrm{g}</code> | g | <code>\mathrm{s}</code> | s | <code>\mathrm{E}</code> | E | <code>\mathrm{Q}</code> | Q |                         |   |
| <code>\mathrm{h}</code> | h | <code>\mathrm{t}</code> | t | <code>\mathrm{F}</code> | F | <code>\mathrm{R}</code> | R |                         |   |
| <code>\mathrm{i}</code> | i | <code>\mathrm{u}</code> | u | <code>\mathrm{G}</code> | G | <code>\mathrm{S}</code> | S |                         |   |
| <code>\mathrm{j}</code> | j | <code>\mathrm{v}</code> | v | <code>\mathrm{H}</code> | H | <code>\mathrm{T}</code> | T |                         |   |
| <code>\mathrm{k}</code> | k | <code>\mathrm{w}</code> | w | <code>\mathrm{I}</code> | I | <code>\mathrm{U}</code> | U |                         |   |
| <code>\mathrm{l}</code> | l | <code>\mathrm{x}</code> | x | <code>\mathrm{J}</code> | J | <code>\mathrm{V}</code> | V |                         |   |

We can also have bold faced Roman English letters using the `\mathbf` command, as in the following table.

|              |              |              |              |              |              |              |              |              |              |
|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|
| $\mathbf{a}$ | $\mathbf{a}$ | $\mathbf{m}$ | $\mathbf{m}$ | $\mathbf{y}$ | $\mathbf{y}$ | $\mathbf{K}$ | $\mathbf{K}$ | $\mathbf{W}$ | $\mathbf{W}$ |
| $\mathbf{b}$ | $\mathbf{b}$ | $\mathbf{n}$ | $\mathbf{n}$ | $\mathbf{z}$ | $\mathbf{z}$ | $\mathbf{L}$ | $\mathbf{L}$ | $\mathbf{X}$ | $\mathbf{X}$ |
| $\mathbf{c}$ | $\mathbf{c}$ | $\mathbf{o}$ | $\mathbf{o}$ | $\mathbf{A}$ | $\mathbf{A}$ | $\mathbf{M}$ | $\mathbf{M}$ | $\mathbf{Y}$ | $\mathbf{Y}$ |
| $\mathbf{d}$ | $\mathbf{d}$ | $\mathbf{p}$ | $\mathbf{p}$ | $\mathbf{B}$ | $\mathbf{B}$ | $\mathbf{N}$ | $\mathbf{N}$ | $\mathbf{Z}$ | $\mathbf{Z}$ |
| $\mathbf{e}$ | $\mathbf{e}$ | $\mathbf{q}$ | $\mathbf{q}$ | $\mathbf{C}$ | $\mathbf{C}$ | $\mathbf{O}$ | $\mathbf{O}$ |              |              |
| $\mathbf{f}$ | $\mathbf{f}$ | $\mathbf{r}$ | $\mathbf{r}$ | $\mathbf{D}$ | $\mathbf{D}$ | $\mathbf{P}$ | $\mathbf{P}$ |              |              |
| $\mathbf{g}$ | $\mathbf{g}$ | $\mathbf{s}$ | $\mathbf{s}$ | $\mathbf{E}$ | $\mathbf{E}$ | $\mathbf{Q}$ | $\mathbf{Q}$ |              |              |
| $\mathbf{h}$ | $\mathbf{h}$ | $\mathbf{t}$ | $\mathbf{t}$ | $\mathbf{F}$ | $\mathbf{F}$ | $\mathbf{R}$ | $\mathbf{R}$ |              |              |
| $\mathbf{i}$ | $\mathbf{i}$ | $\mathbf{u}$ | $\mathbf{u}$ | $\mathbf{G}$ | $\mathbf{G}$ | $\mathbf{S}$ | $\mathbf{S}$ |              |              |
| $\mathbf{j}$ | $\mathbf{j}$ | $\mathbf{v}$ | $\mathbf{v}$ | $\mathbf{H}$ | $\mathbf{H}$ | $\mathbf{T}$ | $\mathbf{T}$ |              |              |
| $\mathbf{k}$ | $\mathbf{k}$ | $\mathbf{w}$ | $\mathbf{w}$ | $\mathbf{I}$ | $\mathbf{I}$ | $\mathbf{U}$ | $\mathbf{U}$ |              |              |
| $\mathbf{l}$ | $\mathbf{l}$ | $\mathbf{x}$ | $\mathbf{x}$ | $\mathbf{J}$ | $\mathbf{J}$ | $\mathbf{V}$ | $\mathbf{V}$ |              |              |

Lowercase Greek symbols are widely used in mathematical formulas and notation. We obtain these symbols using  $\text{\LaTeX}$  commands, which are shown in the following table along with how they appear when rendered.

|            |            |           |           |            |            |
|------------|------------|-----------|-----------|------------|------------|
| $\alpha$   | $\alpha$   | $\iota$   | $\iota$   | $\sigma$   | $\sigma$   |
| $\beta$    | $\beta$    | $\kappa$  | $\kappa$  | $\tau$     | $\tau$     |
| $\gamma$   | $\gamma$   | $\lambda$ | $\lambda$ | $\upsilon$ | $\upsilon$ |
| $\delta$   | $\delta$   | $\mu$     | $\mu$     | $\phi$     | $\phi$     |
| $\epsilon$ | $\epsilon$ | $\nu$     | $\nu$     | $\chi$     | $\chi$     |
| $\zeta$    | $\zeta$    | $\xi$     | $\xi$     | $\psi$     | $\psi$     |
| $\eta$     | $\eta$     | $\pi$     | $\pi$     | $\omega$   | $\omega$   |
| $\theta$   | $\theta$   | $\rho$    | $\rho$    |            |            |

Some of these letters have variants which can be obtained using alternative commands that begin with `var`, as we see in the following table.

|               |               |             |             |             |             |           |           |
|---------------|---------------|-------------|-------------|-------------|-------------|-----------|-----------|
| $\varepsilon$ | $\varepsilon$ | $\varkappa$ | $\varkappa$ | $\varrho$   | $\varrho$   | $\varphi$ | $\varphi$ |
| $\vartheta$   | $\vartheta$   | $\varpi$    | $\varpi$    | $\varsigma$ | $\varsigma$ |           |           |

Given that some uppercase Greek letters are identical to uppercase letters in English, only those that are different to English letters have  $\text{\LaTeX}$  commands or even  $\text{\LaTeX}$  to render them.

These are shown in the following table.

|          |          |           |           |            |            |          |          |
|----------|----------|-----------|-----------|------------|------------|----------|----------|
| $\Gamma$ | $\Gamma$ | $\Lambda$ | $\Lambda$ | $\Sigma$   | $\Sigma$   | $\Psi$   | $\Psi$   |
| $\Delta$ | $\Delta$ | $\Xi$     | $\Xi$     | $\Upsilon$ | $\Upsilon$ | $\Omega$ | $\Omega$ |
| $\Theta$ | $\Theta$ | $\Pi$     | $\Pi$     | $\Phi$     | $\Phi$     |          |          |

A few Hebrew letters are also sometimes used in mathematical notation.

|          |          |         |         |           |           |          |          |
|----------|----------|---------|---------|-----------|-----------|----------|----------|
| $\aleph$ | $\aleph$ | $\beth$ | $\beth$ | $\daleth$ | $\daleth$ | $\gimel$ | $\gimel$ |
|----------|----------|---------|---------|-----------|-----------|----------|----------|

The following calligraphic uppercase English letters are often used in mathematical formulas and notation.

|               |               |               |               |               |               |               |               |               |               |
|---------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|
| $\mathcal{A}$ | $\mathcal{A}$ | $\mathcal{G}$ | $\mathcal{G}$ | $\mathcal{M}$ | $\mathcal{M}$ | $\mathcal{S}$ | $\mathcal{S}$ | $\mathcal{Y}$ | $\mathcal{Y}$ |
| $\mathcal{B}$ | $\mathcal{B}$ | $\mathcal{H}$ | $\mathcal{H}$ | $\mathcal{N}$ | $\mathcal{N}$ | $\mathcal{T}$ | $\mathcal{T}$ | $\mathcal{Z}$ | $\mathcal{Z}$ |
| $\mathcal{C}$ | $\mathcal{C}$ | $\mathcal{I}$ | $\mathcal{I}$ | $\mathcal{O}$ | $\mathcal{O}$ | $\mathcal{U}$ | $\mathcal{U}$ |               |               |
| $\mathcal{D}$ | $\mathcal{D}$ | $\mathcal{J}$ | $\mathcal{J}$ | $\mathcal{P}$ | $\mathcal{P}$ | $\mathcal{V}$ | $\mathcal{V}$ |               |               |
| $\mathcal{E}$ | $\mathcal{E}$ | $\mathcal{K}$ | $\mathcal{K}$ | $\mathcal{Q}$ | $\mathcal{Q}$ | $\mathcal{W}$ | $\mathcal{W}$ |               |               |
| $\mathcal{F}$ | $\mathcal{F}$ | $\mathcal{L}$ | $\mathcal{L}$ | $\mathcal{R}$ | $\mathcal{R}$ | $\mathcal{X}$ | $\mathcal{X}$ |               |               |

Likewise, the following *blackboard bold* fonts are widely used.

|              |              |              |              |              |              |              |              |              |              |
|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|
| $\mathbf{A}$ | $\mathbf{A}$ | $\mathbf{G}$ | $\mathbf{G}$ | $\mathbf{M}$ | $\mathbf{M}$ | $\mathbf{S}$ | $\mathbf{S}$ | $\mathbf{Y}$ | $\mathbf{Y}$ |
| $\mathbf{B}$ | $\mathbf{B}$ | $\mathbf{H}$ | $\mathbf{H}$ | $\mathbf{N}$ | $\mathbf{N}$ | $\mathbf{T}$ | $\mathbf{T}$ | $\mathbf{Z}$ | $\mathbf{Z}$ |
| $\mathbf{C}$ | $\mathbf{C}$ | $\mathbf{I}$ | $\mathbf{I}$ | $\mathbf{O}$ | $\mathbf{O}$ | $\mathbf{U}$ | $\mathbf{U}$ |              |              |
| $\mathbf{D}$ | $\mathbf{D}$ | $\mathbf{J}$ | $\mathbf{J}$ | $\mathbf{P}$ | $\mathbf{P}$ | $\mathbf{V}$ | $\mathbf{V}$ |              |              |
| $\mathbf{E}$ | $\mathbf{E}$ | $\mathbf{K}$ | $\mathbf{K}$ | $\mathbf{Q}$ | $\mathbf{Q}$ | $\mathbf{W}$ | $\mathbf{W}$ |              |              |
| $\mathbf{F}$ | $\mathbf{F}$ | $\mathbf{L}$ | $\mathbf{L}$ | $\mathbf{R}$ | $\mathbf{R}$ | $\mathbf{X}$ | $\mathbf{X}$ |              |              |

As examples of where these fonts are used,  $\mathbb{N}$ ,  $\mathbb{Z}$ ,  $\mathbb{Q}$ ,  $\mathbb{R}$ ,  $\mathbb{C}$  are used to denote the natural numbers, integers, rational, real, and complex numbers, respectively.

## Subscripts and superscripts

Subscripts or superscripts are widely used in mathematical notation. For example, all symbols for variables commonly occur with subscripts and superscripts. In addition, subscripts and superscripts occur in mathematical operators and functions, as we will see below. As we've seen in passing in many examples above, we obtain a subscript by using `_` and a superscript using the `^`. For example, `$x_1$` is  $x_1$  and `$x^2$` is  $x^2$ . When we need to use more than just a single character or command for the subscript or superscript, we need to surround it with braces. For example, if want  $10^{100}$ , we need to write `$10^{\{100\}}`. Writing `$10^100$` will lead to  $10^100$ . Likewise with subscripts. For example, for  $x_{ijk}$ , we write `$x_{\{ijk\}}` and not `$x_ijk$`, which would give  $x_{ijk}$ .

We may also use braces to nest sub and superscripts as in the following example.

`2^{\{2^2\}}`, `2_{\{i_j\}}`

This is rendered as follows:

$$2^{2^2}, 2_{i_j}$$

It should be noted that not using the brace here would simply lead to a ! Double superscript or ! Double subscript error.

It is also possible to mix sub and superscripts. For example, `$2^i_j$` gives  $2_j^i$ . Likewise, `$2_j^i$` also gives  $2_j^i$ .

## Accents and decorations

Symbols for variables sometimes have accents or other decorations. For example, in statistics, it is common to signify a sample mean with a bar over a symbol, e.g.  $\bar{x}$ , or an estimator of a parameter with a hat, e.g.  $\hat{\beta}$ . We obtain these decorations with L<sup>A</sup>T<sub>E</sub>X commands. In the following table are some of the more commonly used accents and decorations.

|                      |           |                       |            |                        |             |                           |                |                        |             |
|----------------------|-----------|-----------------------|------------|------------------------|-------------|---------------------------|----------------|------------------------|-------------|
| <code>\bar{x}</code> | $\bar{x}$ | <code>\dot{x}</code>  | $\dot{x}$  | <code>\vec{x}</code>   | $\vec{x}$   | <code>\mathring{x}</code> | $\mathring{x}$ | <code>\breve{x}</code> | $\breve{x}$ |
| <code>\hat{x}</code> | $\hat{x}$ | <code>\ddot{x}</code> | $\ddot{x}$ | <code>\tilde{x}</code> | $\tilde{x}$ | <code>\acute{x}</code>    | $\acute{x}$    | <code>x^\prime</code>  | $x'$        |

## Arithmetic operations and fractions.

For arithmetic, plus and minus are obtained by `+` and `-`, respectively. For example, `$x + y$` gives  $x + y$ , `$x - y$` gives  $x - y$ , and so on. The command `\pm` is used to produce the *plus or minus* symbol  $\pm$ . For exponents, we use superscripts just as we saw above. For multiplication, it is conventional that the absence of any operator implies multiplication. For example,  $x$  times  $y$  is conventionally often written simply as  $xy$ . However, if we prefer to use a symbol to explicitly mark multiplication, we may use `\cdot` or `\times`. For example, `$a \cdot b$` gives  $a \cdot b$ , and `$a \times b$` gives  $a \times b$ .

For division, if we are simply dividing one symbol by another, we can use `/` or `\div`. For example, `$a / b$` gives  $a/b$ , and `$a \div b$` gives  $a \div b$ . However, for formulas that involve ratios of set of symbols or larger statements, we need to use `\frac{\{ \} \{ \} \}`, as we see in the following example.

`$$`  
`\frac{1}{2} + \frac{3}{4} = \frac{4 + 6}{8} = \frac{5}{4}`  
`$$`

This is rendered as follows:

$$\frac{1}{2} + \frac{3}{4} = \frac{4 + 6}{8} = \frac{5}{4}$$

While it is perhaps more common to use `\frac{}{}` in display mode, it may be used in inline mode as well, as in the following example.

The result is `$x = \frac{a + b}{c + d}$`.

This is rendered as follows:

$$\text{The result is } x = \frac{a+b}{c+d}.$$

We may also nest fractions using `\frac{}{}`, as in the following example.

```
$$
\frac{a}{b + \frac{1}{c}}
$$
```

This is rendered as follows:

$$\frac{a}{b + \frac{1}{c}}$$

### Sums, products, integrals, etc

Summation over multiple variables is denoted using a variant of the uppercase Sigma symbol. However, for this, we use `\sum` and not `\Sigma`. For example, to denote the sum over a set of numbers  $x_1, x_2 \dots x_n$ , we would write this as follows.

```
$$
\sum_{i=1}^n x_i
$$
```

This is rendered as follows:

$$\sum_{i=1}^n x_i$$

The limits of the sum are not strictly necessary. Thus, we could omit them as follows.

```
$$
\sum x_i
$$
```

This appears as follows:

$$\sum x_i$$

However, this summation notation is ambiguous at best, and so including the limits of the sum is highly recommended at all times.

We write products using the uppercase Pi symbol, but again, we should not use `\Pi` but use `\prod` instead. Just like `\sum`, `\prod` should always be used with limits, as in the following example.

```
$$
\prod_{i=1}^n x_i
$$
```

This appears as follows:

$$\prod_{i=1}^n x_i$$

Note that when sums or products are used in inline mode, they appear in a more compact form. For example,  $\sum_{i=1}^n x_i$  appears as  $\sum_{i=1}^n x_i$ , and  $\prod_{i=1}^n x_i$  appears as  $\prod_{i=1}^n x_i$ .

Integrals are written using the `\int` command. Like `\sum` and `\prod`, they may include limits, as in the following example, which gives the area under the curve defined by the function  $f(x)$  between the values of  $x = 0$  and  $x = 1$ .

```
$$
\int_0^1 f(x) dx
$$
```

This appears as follows:

$$\int_0^1 f(x) dx$$

When limits on the integral are given, it is known as a *definite integral*. In the absence of limits, we have an *indefinite integral* and this is the integral over all values of the variable of integrand.

There are, in fact, many other large operators that take limits. In the following table, we show some commonly used ones for set and logical operations.

|                   |                     |
|-------------------|---------------------|
| $\bigcup_{i=1}^n$ | $\bigwedge_{i=1}^n$ |
| $\bigcap_{i=1}^n$ | $\bigvee_{i=1}^n$   |

## Roots

We obtain the square root symbol by the `\sqrt` command. To use this properly, we must include the expression with the `{}` after the command. In other words, to obtain  $\sqrt{2}$ , we must use `\sqrt{2}`. Had we written `\sqrt 2` or even `\sqrt(2)`, we would have obtained  $\sqrt{(2)}$ , which is not the desired result. Note that `\sqrt` produces square root that stretches to enclose the expression to which it applies, as we see in the following example that gives the formula for a sample standard deviation.

```
$$
s = \sqrt{\frac{\sum_{i=1}^n (x_i - \bar{x})^2}{n-1}}
$$
```

This will appear as follows.

$$s = \sqrt{\frac{\sum_{i=1}^n (x_i - \bar{x})^2}{n-1}}$$

To obtain the  $n$ -th root, we place the value of  $n$  in square brackets after `\sqrt` and before the `{}`, as we see in the following example, which displays the cubed root.

```
$$
\sqrt[3]{125} = 5
$$
```

This will appear as follows.

$$\sqrt[3]{125} = 5$$

## Functions

Functions such as trigonometric and logarithmic functions are typeset in Roman font. The following table provides some commonly used examples.



|                      |        |                      |        |                   |     |
|----------------------|--------|----------------------|--------|-------------------|-----|
| <code>\cos</code>    | cos    | <code>\arctan</code> | arctan | <code>\log</code> | log |
| <code>\sin</code>    | sin    | <code>\cosh</code>   | cosh   | <code>\ln</code>  | ln  |
| <code>\tan</code>    | tan    | <code>\sinh</code>   | sinh   |                   |     |
| <code>\arccos</code> | arccos | <code>\tanh</code>   | tanh   |                   |     |
| <code>\arcsin</code> | arcsin | <code>\exp</code>    | exp    |                   |     |

## Equalities, inequalities, set operators

Equality and inequalities can be denoted using `=` and `<` or `>`. For example, `$x = y$` gives  $x = y$ , and `$x < y$` gives  $x < y$ , and so on. However, there are many other commonly symbols used to denote equalities, equivalences, or inequalities. Likewise, there are many commonly used symbols for set theoretic operations. Some of the more widely used examples are shown in the following table.

|                   |                   |                      |                         |                     |                        |                        |
|-------------------|-------------------|----------------------|-------------------------|---------------------|------------------------|------------------------|
| <code>=</code>    | <code>\leq</code> | <code>\gg</code>     | <code>\propto</code>    | <code>\doteq</code> | <code>\subset</code>   | <code>\supseteq</code> |
| <code>&lt;</code> | <code>\geq</code> | <code>\approx</code> | <code>\equiv</code>     | <code>\in</code>    | <code>\supset</code>   | <code>\cap</code>      |
| <code>&gt;</code> | <code>\ll</code>  | <code>\sim</code>    | <code>\triangleq</code> | <code>\ni</code>    | <code>\subseteq</code> | <code>\cup</code>      |

Many of these operators can be negated using the `\not` command before the operator symbol or command. For example, `$x \not= y$` gives  $x \neq y$ , `$x \not< y$` gives  $x \nless y$ , and so. In addition, however, some of these negated operators have their own commands. Likewise, there are some negative operators that are not possible using the `\not`. For example, if we need to indicate that  $x$  is strictly less than  $y$ , we can write  $x \lesseq y$  using `$x \lesseq y$`. Some of these negated operators are shown in the following table.

|                     |                    |                     |                    |                         |                         |
|---------------------|--------------------|---------------------|--------------------|-------------------------|-------------------------|
| <code>\neq</code>   | <code>\ngtr</code> | <code>\ngeq</code>  | <code>\lneq</code> | <code>\nsubseteq</code> | <code>\subsetneq</code> |
| <code>\nless</code> | <code>\nleq</code> | <code>\notin</code> | <code>\gneq</code> | <code>\nsupseteq</code> | <code>\supsetneq</code> |

## Arrows

Arrows are often used in mathematical formulas and notation. The following table provides some of the more commonly used types of arrows.

|                              |                               |                          |
|------------------------------|-------------------------------|--------------------------|
| <code>\leftarrow</code>      | <code>\Longrightarrow</code>  | <code>\Downarrow</code>  |
| <code>\rightarrow</code>     | <code>\leftrightharrow</code> | <code>\nearrow</code>    |
| <code>\to</code>             | <code>\longleftarrow</code>   | <code>\searrow</code>    |
| <code>\longleftarrow</code>  | <code>\Leftrightarrow</code>  | <code>\swarrow</code>    |
| <code>\longrightarrow</code> | <code>\Longleftarrow</code>   | <code>\nwarrow</code>    |
| <code>\Leftarrow</code>      | <code>\uparrow</code>         | <code>\mapsto</code>     |
| <code>\Rightarrow</code>     | <code>\downarrow</code>       | <code>\longmapsto</code> |
| <code>\Longleftarrow</code>  | <code>\Uparrow</code>         |                          |

## Delimiters

Square, round, and angle brackets, curly braces, vertical symbols, slashes, etc., are widely used as delimiters in mathematical expressions. Some of these may be typed using different commands. In the following table, we provide some commonly used delimiters.

|                      |                      |                         |                      |                     |
|----------------------|----------------------|-------------------------|----------------------|---------------------|
| <code>(</code>       | <code>\rbrack</code> | <code>/</code>          | <code>\vert</code>   | <code>\lceil</code> |
| <code>)</code>       | <code>\{</code>      | <code>\backslash</code> | <code>\l</code>      | <code>\rceil</code> |
| <code>[</code>       | <code>\lbrace</code> | <code>\langle</code>    | <code>\Vert</code>   |                     |
| <code>\lbrack</code> | <code>\}</code>      | <code>\rangle</code>    | <code>\lfloor</code> |                     |
| <code>]</code>       | <code>\rbrace</code> | <code> </code>          | <code>\rfloor</code> |                     |

Delimiters that naturally occur in pairs, such as brackets and braces, can be made to stretch to enclose the expressions within them by preceding the left and right forms by `\left` and `\right`, respectively. Consider the following example.

```

$$
\left(1 + \frac{1}{n}\right)^n
$$

```

This appears as follows:

$$\left(1 + \frac{1}{n}\right)^n$$

Had we not used the `\left` and `\right` commands, the brackets would not have stretched, as we see in the following example.

```
$$
(1 + \frac{1}{n})^n
$$
```

This appears as follows, which is not the desired result.

$$(1 + \frac{1}{n})^n$$

Such we wish to stretch delimiters that do not naturally occur in pairs, we may pair them with a `\left.` or `\right.`, as we see in the following example.

```
$$
\left. 1 + \frac{1}{x} \right\vert_a^b
$$
```

This appears as follows.

$$1 + \frac{1}{x} \Big|_a^b$$

We may also use the commands `\big`, `\Big`, `\bigg`, and `\Bigg` to produce increasingly large sizes of the above delimiters, as in the following examples.

```
$$
(x) + \big(x\big) + \Big(x\Big) + \bigg(x\bigg) + \Bigg(x\Bigg)
$$
```

This appears as follows.

$$(x) + (x) + \left(x\right) + \left(x\right) + \left(x\right)$$

Note that these do not stretch to enclose the expressions within them, they are merely different fixed sizes.

## Vectors and matrices

We can produce a matrix using `\begin{matrix}`, `\end{matrix}`, as in the following example.

```
$$
\begin{matrix}
1 & 2 & 3 \\
4 & 5 & 6 \\
7 & 8 & 9
\end{matrix}
$$
```

This appears as follows.

$$\begin{array}{ccc} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{array}$$

Note that here we use `&` to delimit the elements and the `\\` to delimit the rows. There must be an equal number of `&`'s on each row for it to be valid.

Usually, we wish to enclose the matrix with stretching left and right square brackets.

```
$$
\left[
\begin{matrix}
1 & 2 & 3 \\
4 & 5 & 6 \\
7 & 8 & 9
\end{matrix}
\right]
$$
```

This appears as follows.

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix}$$

### Multiline and aligned formulas

We may create multiline formulas that are vertically aligned at designated points using `\begin{aligned}`, `\end{aligned}`, as in the following example.

```
$$
\begin{aligned}
f &= (x + y)(x + y), \\
&= x^2 + 2xy + y^2
\end{aligned}
$$
```

This appears as follows.

$$\begin{aligned} f &= (x + y)(x + y), \\ &= x^2 + 2xy + y^2 \end{aligned}$$

Note that similarly to the case of the `matrix` environment, we use `&` for alignment and `\\` for newlines.

We may use the `aligned` environment even when we do not need alignment per se, but just require multiline equations, as in the following example.

```
$$
\begin{aligned}
x &= a + b \\
y &= c + d + e
\end{aligned}
$$
```

This appears as follows.

$$x = a + b$$

$$y = c + d + e$$

### Spacing in math mode

L<sup>A</sup>T<sub>E</sub>X will usually take care of the spacing between all elements of a formula. Occasionally, however, we need to adjust space. To do so, we can use the commands `\,`, `\:`, `\;`, `\` (i.e., a `\` followed by a space; this gives an interword space), `\quad` and `\qquad` to produce increasing large spaces. As we see in the following table, blank spaces on their own do not introduce any additional spaces.

|                     |                 |                         |                  |
|---------------------|-----------------|-------------------------|------------------|
| <code>xy</code>     | <code>xy</code> | <code>x \quad y</code>  | <code>x y</code> |
| <code>x y</code>    | <code>xy</code> | <code>x \qquad y</code> | <code>x y</code> |
| <code>x \, y</code> | <code>xy</code> |                         |                  |
| <code>x \: y</code> | <code>xy</code> |                         |                  |
| <code>x \; y</code> | <code>xy</code> |                         |                  |
| <code>x \ y</code>  | <code>xy</code> |                         |                  |

### Text in math mode

If we require normal English text within a formula, particularly within display mode, it is necessary to enclose the text using the `\text` command, as in the following example.

```
$$
x = y, \quad \text{this is text}
$$
```

This appears as follows.

$$x = y, \quad \text{this is text}$$

Had we placed the English text in the formulas without using `\text`, we would not obtain what we want.

```
$$
x = y, \quad \text{this is text}
$$
```

This appears as follows, which is not the desired outcome.

$$x = y, \quad \textit{this is text}$$

Often when we use text in displayed formulas, we wish to include mathematical symbols and notations inline in the text. To do so, we use inline math mode within the text within the `\text` command, as in the following example. Had we placed the English text in the formulas without using `\text`, we would not obtain what we want.

```
$$
s_i = 2^i, \quad \text{for } i \geq 0
$$
```

This appears as follows.

$$s_i = 2^i, \quad \text{for } i \geq 0$$

## Underbraces, overbraces

We may use `\underbrace` and `\overbrace` to produce horizontally curly braces. These can be used to annotate parts of a formula, as in the following example.

```
$$
\underbrace{(x + y)(x - y)}_{\text{left hand side}}
= \overbrace{x^2 - y^2}^{\text{right hand side}}
$$
```

This appears as follows.

$$\underbrace{(x + y)(x - y)}_{\text{left hand side}} = \overbrace{x^2 - y^2}^{\text{right hand side}}$$

Note that we must use `\text` if we wish to include normal text in the annotation. Also, we use `_` and `^` after `\underbrace` and `\overbrace`, respectively, for the annotation.

## Miscellaneous symbols

There are many more symbols available for use in math modes in Markdown. The following table provides a few commonly used symbols that were not mentioned above.

|                       |            |                      |           |                          |               |                     |         |
|-----------------------|------------|----------------------|-----------|--------------------------|---------------|---------------------|---------|
| <code>\infty</code>   | $\infty$   | <code>\exists</code> | $\exists$ | <code>\emptyset</code>   | $\emptyset$   | <code>\cdots</code> | $\dots$ |
| <code>\partial</code> | $\partial$ | <code>\top</code>    | $\top$    | <code>\varnothing</code> | $\varnothing$ |                     |         |
| <code>\nabla</code>   | $\nabla$   | <code>\dagger</code> | $\dagger$ | <code>\Re</code>         | $\Re$         |                     |         |
| <code>\forall</code>  | $\forall$  | <code>\dagger</code> | $\dagger$ | <code>\ldots</code>      | $\dots$       |                     |         |

## Git

Version control software (VCS) is an essential tool for the efficient management and organization of source code. In the case of data analysis using R, the relevant source code files will primarily include `.R` and `.Rmd` scripts, but even in small and routine projects, there are many other possibilities too. VCS allows us to keep track of all the versions or revisions to a set of files in an efficient and orderly manner. As a simple example to motivate the use of a VCS system, let us say that we are working on a relatively small data analysis project initially involving some `.R` and `.Rmd` scripts, with names like `preprocessing.R`, `exploration.R`, `analysis.R`, and `report.Rmd`. Let's say that we work with these files, adding new code, editing or deleting old code, etc., every few days in a normal R session. If we were to simply save the files after each session, we would obviously only ever have their most recent versions. All the previous versions would be lost. In order to avoid loss of previous versions in case they are needed, we could periodically *save as*, creating versions like `preprocessing_v1.R`, `analysis_oct23.R`, and so on. As time goes by, it is highly likely too that new files will appear. Some of these may have been only intended to be temporary files, but others might be intended to be vital parts of the project. Some new file might be *branches* of other files, where we copy the original, and work on some new feature of the code in the copy with the intention of merging the changes back if and when necessary. By proceeding in this manner, there is usually an eventual proliferation of new files and different versions of files with sometimes ambiguous or inscrutable names like `analysis_v1_tmp.R`, `analysis_v1_tmp_new.R`, `preprocessing_foo_tmp_foo.R`, and so on. If files are being copied between different devices or to cloud based storage, and edited on different devices, the situation can get ever more disorganized, with files of similar names but perhaps slightly different contents or different time-stamps across different machines. At this point, especially if we return to this work after a period of time, it is not usually not clear even what each file does, where the latest version of any file is, not to mention what all the previous or temporary versions contain and when and why they were made. If we collaborate with others, things usually become even worse. First, we must decide on a means of sharing files, with email attachments still probably being the default and most widely used method of doing so. Sending back and forth emails with modifications creates yet more versions to manage, and multiple people working independently the same files introduces conflicts that need

to be manually resolved. Eventually, we have multiple files and versions, on multiple devices, being edited independently by multiple different people. Knowing what each file and version is or does, and who did what and when and where is usually lost as a result.

This level of disorganization is frustrating, wasteful of time and effort, and obviously bad for reproducibility. The authors themselves may find it difficult or impossible to pick through their files to recover and reproduce all the steps involved in any analyses. Moreover, even if they were only working on one file such as an RMarkdown file from which their final report was generated, the proliferation of versions across different devices and owners would still occur, making it difficult to pick up and resume their work after a period of inactivity. In addition, they may lose track of which version of the `.Rmd` produced which version of the rendered manuscript. It is all very well knowing that a manuscript was produced by knitting a `.Rmd`, and hence that all its reported results are reproducible in principle, but if we have lost track of the `.Rmd` that produced it, it is obviously no longer reproducible in practice.

VCS systems allow us to manage our source files in an orderly and efficient manner. There are many VCS systems available, both proprietary and open source, and while precise information on usage worldwide is hard to establish definitively, almost all surveys of VCS usage show that *Git* is now by far the most popular and widely used VCS system. Git is open-source software was originally developed in 2005 for version control of the development of the Linux operating system, something for which it is still used. It gradually became more widely used in the open source community and within a few years had become popular than the previously very widely used open source *subversion* VCS system. With the growing popularity of Git hosting sites like GitHub, which currently hosts over 100 million Git based projects, Git is now the most widely used VCS system worldwide.

In what follows, we aim to provide a tutorial introduction to the main features of Git, the main Git commands, and how to use it. Obviously, it is beyond the scope of this section to provide a comprehensive introduction. This can be found in many books such as Chacon and Straub (2014), which is available in its entirety online at <https://git-scm.com/book/en/v2>.

## Installation

Git is available for Windows, MacOS, Linux. Git is first and foremost a command driven software. There are graphical interfaces, i.e. GUIs, to Git, but we will not consider them here and do not recommend them either given that there is only a small number of core commands to learn and they allow Git to be used both efficiently and identically across all different devices.

For Windows, we highly recommend installing and using the *Git Bash* shell available from <https://gitforwindows.org/>. This provides a Bash Unix shell<sup>1</sup> from which Git can be used just as it would be used on other Unix systems like Linux<sup>2</sup> and MacOS.

For MacOS, Git is already preinstalled on recent versions of MacOS. While this may be perfectly adequately, the preinstalled Git on MacOS is based on a build of Git by Apple and is usually not up to date with the latest version of Git. More recent versions of Git for MacOS are available elsewhere, such as <https://git-scm.com/download/mac>.

For Linux, given the role of Git in the development of Linux, Git is seen as a vital Linux tool. It is easily installed using the package managers of any Linux distribution, see <https://git-scm.com/download/linux>.

Once Git is installed, it is available for use using the command `git` in an operating system terminal. For Windows, this means that it is available in the DOS shell using the command `git`. However, if Git is installed, as we recommend, as the Git Bash shell, that Bash shell should be always used instead of the DOS shell. For MacOS and Linux, Git will be available in the system terminal, and will work identically in the Bash, sh, zsh, etc., shells.

---

<sup>1</sup>[https://en.wikipedia.org/wiki/Bash\\_\(Unix\\_shell\)](https://en.wikipedia.org/wiki/Bash_(Unix_shell))

<sup>2</sup>Strictly speaking, Linux is not Unix, but rather a \*nix or Unix-like operating system. However, it is essentially a free and open-source re-implementation of Unix, and so can be seen as Unix for all practical purposes.

For what follows, we will assume users are using a Unix shell, and so all of this will be equally applicable to users of Linux, MacOS, and Windows users assuming they use Git Bash.

To establish that Git has been successfully installed and is available for use, type the following

```
git --version
```

The output of this command on a relatively up to date (as of May, 2020) version could appear as follows.

```
git version 2.26.2
```

## Configuration

Before we start using Git, we need to perform some minimal configuration. Specifically, we first need to set our name and email address. Git requires that each time we *commit* to the repository, as we will see below, we have the name and email address of the person doing the committing. This information could be set on a per commit basis. However, it is more common and easier to set this information as a global configuration setting and then it will be used whenever the user performs a commit. This can be done using `git config` as in the following code.

```
git config --global user.name 'Mark Andrews'
git config --global user.email 'mjandrews.org@gmail.com'
```

It is recommended that we also need to set the text editor that we will use for writing our commit messages, which is also something we will see below. By default, the text editor is the `vi` or `vim` editor. These editors are standard Unix editors. They are loved by some and loathed by others. To the uninitiated, these editors are likely to be seen as difficult and probably annoying to use. In any case, it is certainly not necessary to use them as any text editor can be used instead. If a user already has a preference, they should use this. If not, one recommended editor, which is open-source editor and available across the Windows, MacOS, and Linux platforms, is the *Atom* editor, see <https://atom.io/>. Atom can be set as the default editor to be used with Git as follows.

```
git config --global core.editor "atom --wait"
```

Once set, this configuration information is stored in a *dot file* named `.gitconfig` in the user's home directory. Dot files are a standard Unix file system feature. They are simply files, or possibly even directories, that begin with a dot. They are hidden by default in file listings, and are primarily intended to be used for configuration information. The `.gitconfig` file can be edited at any time to change global configuration settings.

## Creating and initiating a Git repository

A Git repository is simply a directory (i.e., a folder) in which all the files, including those in subdirectories, are being tracked and managed, or potentially tracked and managed, by Git. Sometimes, as we will see, we obtain a Git repository by *cloning* it. However, we can always turn any directory on our computer into a Git repository using a single command, as we will see momentarily.

Before we proceed a few comments and warnings are in order. First, repositories are independent and self-contained entities. By independent, we mean that separate Git repositories do not depend or interact with one another. There are in fact some exceptions to this, namely by using Git submodules and subtrees. However, these are used in only a small proportion of cases and we will not consider them further here. By self-contained, we mean that each repository contains all its own content tracking and history information. In other words, the vital information for using the repository is contained within the repository itself and not outside of it in, for example, a system wide database or on a server. Now, as a word of warning, although you can turn any directory a Git repository, we should not have Git repositories *within* Git repositories. For example, if there is a directory `foo` inside and another directory named `foobar`, if we turn `foo` into a Git repository, we should not do so to `foobar` as well. Likewise, if we make `foobar` a Git repository, we should not make `foo` a Git repository too. In general then, any ancestor directories, e.g. parent, grandparent,

directories, or any descendant directories, e.g. child, grandchild directories, of any Git repository should not itself also be a Git repository.

Now let us create a Git repository from a newly created empty directory. There is no necessity for the directory to be newly created and empty, that's just what we use this example. First, let us create a new empty directory `project101` inside of the `Documents` directory in our home directory. We will assume that the directory `Documents` already exists in our directory, but this is quite common across different platforms. We will use the standard Unix commands `mkdir` to make the `project101` directory, and then use the Unix command `cd` to change directory into it. Also, on Unix systems, the `~` means the user's home directory. Note that here, the `mkdir` command would raise an error if a directory called `project101` already exists.

```
mkdir ~/Documents/project101
cd ~/Documents/project101
```

We may list the contents of `project101` using the Unix `ls` command. Here, we use `-laF` option to `ls`. The `-l` asks for the information to be shown with one file or directory per line. The `-a` option will show so-called "hidden" files and directories. Hidden files and directories, also known as *dot* files or directories, begin with `..`. They are intended to contain configuration information and, by default, not shown in lists of files and directories. The `-F` option is used primarily to indicate whether the items in a directory are files or subdirectories. Directories are listed by ending their name with a `'/'`.

```
ls -laF
```

```
./
../
```

From this listing, we see that the directory is empty. The two items that are listed, `./` and `../`, are always present in a listing where we show hidden files. They are merely references to the present (`.`) and parent (`..`) directory.

Now, we create a Git repository in `project101` as follows.

```
git init
```

```
Initialized empty Git repository in /home/andrews/Documents/project101/.git/
```

From this, we see that Git initializes a repository in `project101`, and created a directory therein named `.git`. We may now look at the contents of `project101` again.

```
ls -laF
```

```
./
../
.git/
```

We see now that there is a hidden or *dot* directory named `.git`. The presence of the `.git` subdirectory is a necessary and sufficient condition for the directory to be Git repository.

We can now run the `git status` command to see what is the current state of the repository.

```
On branch master
```

```
No commits yet
```

```
nothing to commit (create/copy files and use "git add" to track)
```

From the output of this command, we see a number of important pieces of information. First, it tells us that we are on the `master` branch of the repository. What branches are and how to use them is something to which we will return in due course below. Suffice it to say for now that they are a major feature of Git repositories. Next, we see that there are `No commits yet`. This means that we have not yet *committed* anything to the repository and so nothing is being tracked or managed by Git yet. Finally, we also see that there is currently



nothing to commit. In other words, there are no files in `project101` that could potentially be committed to the repository.

Now, let us put some files into `project101`. These files could be any type of file, but because Git is intended for the management of source code, ideally the files should be text files rather than binary files. For this example, we will put one R script, `script.R`, and one other text file, `readme.md`, into `project101`. These files could be simply moved or copied from some other directory or could be written in an editor, such as RStudio, and then saved into `project101`. For this example, we will write these files using the Unix `cat` command and file redirection methods, which can always be used to write lines of text to files.

```
cat << EOF > script.R
library(tidyverse)

survivors <- Titanic %>%
  apply(c('Sex', 'Survived'), sum)
EOF

cat << "EOF" > readme.md
The project contains an R script, `script.R`, for processing the `Titanic` data set.
EOF
```

With this, we'll create two files, `script.R` and `readme.md`. As we can see, `script.R` has a few simple lines of R code and `readme.md` has one line of text explaining what the contents of `project101` does.

After we have put these files in `project101`, we can do a file listing as we did above.

```
ls -laF
```

```
./
../
.git/
readme.md
script.R
```

We see that `script.R` and `readme.md` are there. Next, we can do a `git status` as above.

```
git status
```

```
On branch master
```

```
No commits yet
```

```
Untracked files:
```

```
(use "git add <file>..." to include in what will be committed)
  readme.md
  script.R
```

nothing added to commit but untracked files present (use "git add" to track)

Much of the information here is as it was previously, but now we are told that there are two **Untracked files**, namely `readme.md` and `script.R`. Files listed as *untracked* mean that they are present in the directory but not as yet being tracked or managed by Git. As the output indicates, however, we can use the command `git add` to get Git to track them, as we do in the following command.

```
git add readme.md script.R
```

Now, let us check the status of the repository again.

```
git status
```

On branch master

No commits yet

Changes to be committed:

(use "git rm --cached <file>..." to unstage)

new file: readme.md

new file: script.R

We see from the output that there are new files, `readme.md` and `script.R`, that *can* be committed. This is an important point. The files are not yet committed to the repository, they are *staged* for commitment. When files are staged, they are in an intermediate area, kind of like a departure lounge in an airport. To commit them, we must run the `git commit` command as follows.

This will open your editor, and it will contain the following text exactly.

```
# Please enter the commit message for your changes. Lines starting
# with '#' will be ignored, and an empty message aborts the commit.
#
# On branch master
#
# Initial commit
#
# Changes to be committed:
#   new file:   readme.md
#   new file:   script.R
#
```

In other words, the lines beginning with `#` provide information to you as you write your commit message, but they will not be part of the message itself. You write your message above these `#` lines. It is conventional and recommended<sup>3</sup> that the first line of this message is no more than 50 characters long and is followed by a blank line and then followed by more elaboration. The first line is treated as a subject line. Its first character should be capitalized, and it should not end in a full-stop/period. It is also recommended that this subject line be written in imperative tense and not past tense. For example, it is recommended that we write something like “Add new function ...” rather than “Added new function ...”. Admittedly, this is probably initially an unnatural way to write for most people. After the subject line, there must be a blank line. Without it, some features of Git can be affected. Then, a more elaborate message can be written. Here, the imperative tense is no longer necessary. In fact, it is not necessary to have a body at all, but it is highly to provide a message body and to use it provide details about what the code being committed does. It will be helpful for others, including your future self, to understand what was being added and why. The character line of the body text is recommended to to 72. Most Git-aware editors, like Atom, will indicate if the subject line is over 50 characters, and if there is not a blank line after the subject, and will wrap the body at 72 characters.

In this example, because this commit is our first commit, our subject line should acknowledge that this is the beginning of the project. As such, a message like the following can be used.

Initialize the repository

Two files are added.

- \* ``script.R`` is an R script summarizing the ``Titanic`` data-set
- \* ``readme.md`` is the project's readme file.

---

<sup>3</sup>See <https://tbagery.com/2008/04/19/a-note-about-git-commit-messages.html>

Let us assume, therefore, that we this is what we have typed into the editor after we ran the `git commit` command the editor was opened. After we save this file and quit the editor, the commit is completed.

Before proceeding, we should note that it is also possible to write one-liner commit messages using the option `-m` as follows.

```
git commit -m 'Do something to something'
```

Here, no editor opens and there's no body to the message. For simplicity, we will use this method below sometimes. However, using this method is probably to be avoided as it will, by necessity, lead to minimal and probably poorly thought out commit messages.

We can now view at the Git log of the repository with the following command.

```
git log
```

The output of `git log` would be something like this.

```
commit 4cb519c9dc0579eb285c1f00ed6d2163dbcda1ef
Author: Mark Andrews <mjandrews.org@gmail.com>
Date:   Sun May 24 18:07:30 2020 +0100
```

Initialize the repository

Two files are added.

- \* ``script.R`` is an R script summarizing the ``Titanic`` data-set
- \* ``readme.md`` is the project's readme file.

As we can see, our commit message is there, as is our name, email address, and the date and time stamp of the commit. A crucial additional piece of information is the *commit hash*, which in this case is

```
4cb519c9dc0579eb285c1f00ed6d2163dbcda1ef
```

The commit hash is a vital feature of Git. It is a 40 hexadecimal character (160 bit) cryptographic hash of the contents and other defining information about each commit. In other words, it can be seen as essentially a fingerprint of the commit, but not just an arbitrarily assigned identifier, but one that is calculated using a hashing algorithmic from the contents of the commit. It can therefore be used to uniquely identify the commit and to do a file integrity check of its contents.

Having performed the commit, we can now again check the status of the repository.

```
git status
```

```
On branch master
```

```
nothing to commit, working tree clean
```

As we can see, now the repository is in a clean state. All the original files have been committed. There are as yet no new files in the directory, and no edits to the existing files yet either.

As an interim summary, thus far, we have seen a number of essential and regularly used Git commands.

**git init** Create a Git repository in the current working directory of the shell.

**git status** Report the current working status of the repository. Specifically, are there untracked files in the repository or files in the staging area that have not yet been committed.

**git add <file> ...** Add the files <file> ... to the staging area. The staging area is like the departure lounge of an airport. The files therein are going scheduled for committal, though they may be taken out of the staging area too.

**git commit** Commit the files in the staging area. This command opens an editor and a commit message is entered there.

## Adding and editing files

As with any data analysis project, as it progresses, new files will be made and edits will be made to existing one. With Git, we can choose whether and when to add new files to the repository. In other words, for example, if new files are added to the directory `project101`, they will be treated as *untracked* files. They will never be automatically added to the repository. We must explicitly add them, using `git add`. Moreover, as we've seen above, adding files with `git add` only puts them in the staging area for committal. They are not committed until we explicitly commit them with `git commit`. Something similar occurs with edits to the existing files. After any edits, Git identifies that files have be modified. However, for these changes to be committed to the repository, they must first be added to the staging area with the `git add` command. Then, they must be explicitly committed with a `git commit` command. This double step process allows us to build up the staging area gradually as we work, and then committing all its contents when we they are all ready. This intended purpose of this is that we can then commit a set of files and edits that are all related to one another and together effectively do one main thing, such as fix a bug or a new feature, etc. These are known as *atomic commits*.

Continuing with our very simple example project, let us now add a new file. As above, for illustrative purposes here, though this is not likely to be used in practice, we will use the Unix `cat` and file redirection commands to accomplish this.

```
cat << "EOF" > models.R
model <- lm(Fertility ~ Catholic, data = swiss)
model_summary <- summary(model)
```

```
bash: line 2: warning: here-document at line 0 delimited by end-of-file (wanted `EOF')
```

Now, let us edit the `script.R` file. In practice, we would most likely edit `script.R`, but for simplicity here, we will use another Unix command `sed` (*stream editor*) to edit some of the file's contents. In particular, we will use `sed` to change the function `sum` for `mean` in the `apply` function.

```
sed -i 's/sum/mean/' script.R
```

Here, `s/sum/mean` means substitute `mean` for `sum` on every line. In this example, `sum` only occurs on one line, namely where the `apply` function is used, and so this command replaces `mean` for `sum` as the function in `apply`. The `-i` option means *in place*. In other words, do not print to standard output the lines of `script.R` with the changes applied, but apply the changes to `script.R`.

Now, let us check the working state of the repository.

```
git status
```

```
On branch master
```

```
Changes not staged for commit:
```

```
  (use "git add <file>..." to update what will be committed)
```

```
  (use "git restore <file>..." to discard changes in working directory)
```

```
    modified:   script.R
```

```
Untracked files:
```

```
  (use "git add <file>..." to include in what will be committed)
```

```
    models.R
```

```
no changes added to commit (use "git add" and/or "git commit -a")
```

From this, we see that there is, as expected, one untracked file, `models.R`, and one modified file, `script.R`.

We may now add both of these files, and the commit them both to the repository at the same time using a single commit. Alternatively, we could add and commit them individually. Whether we proceed one way or another should be based on whether the commit is atomic, i.e. has one main unitary function or purpose. In this example, because the new file and the change to `script.R` are not related to one another, we will

perform two separate commits. First, we will add the modified `script.R` to staging.

```
git add script.R
```

Before we proceed, let us check the working state of the repository.

```
git status
```

```
On branch master
```

```
Changes to be committed:
```

```
  (use "git restore --staged <file>..." to unstage)
```

```
    modified:   script.R
```

```
Untracked files:
```

```
  (use "git add <file>..." to include in what will be committed)
```

```
    models.R
```

As we can see, `script.R` is now in the staging area ready to be committed. We now can do the commit.

```
git commit
```

This will bring up our editor with the following contents.

```
# Please enter the commit message for your changes. Lines starting
# with '#' will be ignored, and an empty message aborts the commit.
#
# On branch master
# Changes to be committed:
#   modified:   script.R
#
# Untracked files:
#   models.R
#
```

As before, we should write our commit message above the lines beginning with `#`, following the conventions and recommendations for good git messages. We will add the following text to the commit message in the editor.

```
Change function used in the apply functional
```

```
The functional now calculates the mean number of men and
women who survived the Titanic, not the total number.
```

Then, as before, when we save and quit, the staged changes are committed. If we look at the logs, we will now see the following.

```
git log
```

```
commit b878a06b98841d661952ca1c0acdd826bee717a3
Author: Mark Andrews <mjandrews.org@gmail.com>
Date:   Sun May 24 19:07:30 2020 +0100
```

```
    Change function used in the apply functional
```

```
    The functional now calculates the mean number of men and
    women who survived the Titanic, not the total number.
```

```
commit 4cb519c9dc0579eb285c1f00ed6d2163dbcda1ef
Author: Mark Andrews <mjandrews.org@gmail.com>
```

Date: Sun May 24 18:07:30 2020 +0100

Initialize the repository

Two files are added.

- \* `script.R` is an R script summarizing the `Titanic` data-set
- \* `readme.md` is the project's readme file.

Let us again check the repository's status.

On branch master

Untracked files:

(use "git add <file>..." to include in what will be committed)  
models.R

nothing added to commit but untracked files present (use "git add" to track)

As expected, we now see that there is just one untracked file, `models.R`, in the directory. We first add this file to put it in the staging area.

```
git add models.R
```

Then, we commit it with the following message.

Add new lm script named `models.R`

This script performs a linear regression analysis of the built-in `swiss` data-set.

The status of the repository should confirm that everything is now clean.

```
git status
```

On branch master

nothing to commit, working tree clean

We can also check the logs.

```
git log
```

```
commit 5dafc885149f08a35fbb7b1a43668a07ea303a2f
Author: Mark Andrews <mjandrews.org@gmail.com>
Date: Sun May 24 19:07:30 2020 +0100
```

Add new lm script named `models.R`

This script performs a linear regression analysis of the built-in `swiss` data-set.

```
commit b878a06b98841d661952ca1c0acdd826bee717a3
Author: Mark Andrews <mjandrews.org@gmail.com>
Date: Sun May 24 19:07:30 2020 +0100
```

Change function used in the apply functional

The functional now calculates the mean number of men and women who survived the Titanic, not the total number.

```
commit 4cb519c9dc0579eb285c1f00ed6d2163dbcda1ef
```

Author: Mark Andrews <mjandrews.org@gmail.com>  
Date: Sun May 24 18:07:30 2020 +0100

Initialize the repository

Two files are added.

- \* ``script.R`` is an R script summarizing the ``Titanic`` data-set
- \* ``readme.md`` is the project's readme file.

As expected, we now have three commits in the log now.

## Reverting changes

With Git, as with any VCS system, we can effectively undo any commits we have made or revert back to any previous state in the repository. We may also simply view or copy any file from at any point in the past, or we may temporarily move the repository back to any of its previous states. The principal Git commands that we use to accomplish these and related actions are `git checkout`, `git show`, `git reset`, and `git revert`. With all of these commands, we usually use the commit hash as the unique identifier of the repository's state. We have already seen that each hash is a 40 hexadecimal character string, and they are listed in the repository's logs. If want to view just the commit hash in the logs, we can do so with the following command.

```
git log --pretty=format:@"%H"
```

```
5dafc885149f08a35fbb7b1a43668a07ea303a2f
b878a06b98841d661952ca1c0acdd826bee717a3
4cb519c9dc0579eb285c1f00ed6d2163dbcda1ef
```

We can also view the short form of the hash, which is the first 7 characters, as follows.

```
git log --pretty=format:@"%h"
```

```
5dafc88
b878a06
4cb519c
```

The short form of the hash is usually, i.e. with very high probability, sufficient to uniquely identify it, and so it is very common practice to use the short form as the commit identifier in Git commands.

## Git checkout

With the `git checkout` command, we can return to the state of the repository at any point in history. As an example, let us return to the state of the repository after the second commit. The short form of the hash of this commit is `b878a06`. We can use the `git checkout` command to return to this hash as follows.

```
git checkout b878a06
```

Doing so, by default, we will by get output along the following lines.

Note: switching to 'b878a06'.

You are in 'detached HEAD' state. You can look around, make experimental changes and commit them, and you can discard any commits you make in this state without impacting any branches by switching back to a branch.

If you want to create a new branch to retain commits you create, you may do so (now or later) by using `-c` with the switch command. Example:

```
git switch -c <new-branch-name>
```

Or undo this operation with:

```
git switch -
```

Turn off this advice by setting config variable `advice.detachedHead` to `false`

HEAD is now at b878a06 Change function used in the apply functional

This is telling us that we are in a `detached HEAD` state. This requires some brief explanation.

If we have a Git repository with a single branch, known as the *master branch* by default, as has been the case so far in this example, we can view each commit as a point on a chain. The tip of the chain is known as the `HEAD`. Or more precisely, the `HEAD` is a pointer to the tip of the chain. We can visualize this as follows, where each `o` is a state of the repository.

```
master branch: o -- o -- o -- o --- o <-- tip
                |
                HEAD
```

New commits on the master branch must go on top of the tip of this branch. If we do a `git checkout <commit hash>`, as we did above, the `HEAD` now points to some intermediate point in the chain. The master branch, and its tip, is unaffected.

```
master branch : o -- o -- o -- o --- o <-- tip
                |
                HEAD
```

As the message says we can look around, make changes and commit at this point, start a new branch here, etc., but this won't affect the master branch. As such, probably the most common use of the `git checkout <commit hash>` command is simply to jump back to some point in the repository's history and look around at the files and their contents. This is very useful. However, the important point is to remember that we have not actually rolled the repository's current state back to that point yet.

Having gone back to `b878a06`, we can, for example, do a file listing.

```
ls -laF
```

```
./
../
.git/
readme.md
script.R
```

As we see, we have just the `script.R` and `readme.md` files. The `models.R` file is not there because this was added at the next point in the repository's history. Likewise, we look at the log, we see only two commits.

```
git log --pretty=format:"%h"
```

```
b878a06
4cb519c
```

In other words, we see the history of the repository as it was at commit state `b878a06`.

Now, having looked around at `b878a06`, we can return the current state of the repository, i.e. the tip of the master branch. We do so as follows.

```
git checkout master
```

Previous HEAD position was b878a06 Change function used in the apply functional  
Switched to branch 'master'



## Git show

If we just wanted to view a single file as it was in some previous point in the repository's history, there is an alternative command to use: `git show`. For example, we want to view the file `script.R` as it was on commit `4cb519c`, we could do as follows.

```
git show 4cb519c:script.R
```

```
library(tidyverse)
```

```
survivors <- Titanic %>%  
  apply(c('Sex', 'Survived'), sum)
```

As we can see, this is exactly as `script.R` was on the initial commit. The contents of the file we `show` is output to the standard output. If we want to copy these contents to a new file, e.g., named `script_previous.R`, we can do so with standard Unix file redirections, as in the following example.

```
git show 4cb519c:script.R > script_previous.R
```

Now, a standard file listing will show that this new file is present in the working directory.

```
ls -laF
```

```
./  
../  
.git/  
models.R  
readme.md  
script_previous.R  
script.R
```

Because we will not use it further in this example, we will remove it now.

```
rm script_previous.R
```

## Git reset

Git's `reset` command can be seen as related to its `checkout` command, but it moves both the head and tip of the branch back to the designated commit. For example, let us image this is the current state of the master branch and its history.

```
master branch: o -- o -- o -- o --- o <-- tip  
                |  
                HEAD
```

If we `reset` to a previous commit, the tip of the master branch is assigned to there.

```
master branch: o -- o -- o <-- tip  
                |  
                HEAD
```

In this example, the two commits that were reset are not immediately lost, but will be removed as soon as the next Git *garbage collection* is done. Garbage collection can be done manually at any time using `git gc`, and some Git commands run it automatically.

Let's do a `reset` to `b878a06`.

```
git reset b878a06
```

Now, let us look at the status.

```
git status
```

On branch master

Untracked files:

```
(use "git add <file>..." to include in what will be committed)
models.R
```

nothing added to commit but untracked files present (use "git add" to track)

We see that the file `models.R`, which was what was committed in the last commit, is in the working directory. In other words, we have returned to the state of the repository just prior to the `git add` command where we added `models.R` to the staging area.

There are three main *modes* to `git reset`: *soft*, *hard*, and *mixed*. The *mixed* mode is the default, which we have just used, but can also be obtained explicitly with the `--mixed` option. The defining feature of *mixed* mode is that the changes that were made in the reset commits are not yet staged, but they are present in the working directory, either as untracked files or as modifications to existing files.

In our case, because we used the default *mixed* mode, the file `models.R` is present as an untracked file in the working directory. If we effectively want to undo our resetting, we can add this file, as per usual.

```
git add models.R
```

Then, we can commit this, again as per usual. Here, we will in fact use the exact same commit message we used previously, whose subject line was as follows.

```
Add new lm script named `models.R`
```

We will add this text to the file that opens when we run the following command.

```
git commit
```

Now, let us look at the logs.

To keep it brief, in this case, we will just show the short form of the hash and the subject line.

```
git log --pretty=format: '%h %s'
```

```
7f37a5f Add new lm script named `models.R`
b878a06 Change function used in the apply functional
4cb519c Initialize the repository
```

As we can see, the commit of the most recent change has changed now. Previously, it was `5dafc88`, and now it is `7f37a5f`. This is because the date, amongst other things, is used to calculate the commit hash.

In *soft* reset mode, we again reset to the a previous commit. This time, we changes that we reset are present in the staging area. In other words, we reset to the point *after* the adding to the staging area was done, but prior to the committing. In the present example, we can do a soft reset to the penultimate commit as follows.

```
git reset b878a06 --soft
```

Now, let us look at the status.

```
git status
```

On branch master

Changes to be committed:

```
(use "git restore --staged <file>..." to unstage)
new file:   models.R
```

As we can see, in this case, the `models.R` is in the staging area. Now, we can effectively undo our reset by performing a commit

```
git commit
```

If we enter the same commit message here as we used originally, we will effectively return to the original tip of the master, albeit with a new hash and new date.

The final reset mode to consider is the *hard*. This will not only undo the relevant commits but it will also unstage anything in the staging area, and remove any untracked or modified files from the directory. As it will destroy permanently destroy your changes, it ought to be used with caution. If we wished to perform a hard reset in this example, which we will not actually do here, the command would be as follows:

```
# git reset b878a06 --hard
```

## Git revert

A final command for effectively undoing changes and rolling back history is `git revert`. Unlike `git reset`, `git revert` leaves the previous history and state of the repository intact. It adds new changes to the repository to effectively recreate its contents at a previous points.

Using our example, let us **revert** back to the first commit in the repository. As a reminder, the set of commits in our history is the following.

```
git log --pretty=format: '%h %s'
```

```
37790d0 Add new lm script named `models.R`
b878a06 Change function used in the apply functional
4cb519c Initialize the repository
```

As such, we want to revert back to 4cb519c. This means that we want to undo the change applied at 37790d0 and then undo the change applied at b878a06. There are alternative syntax options to do this, but we will use the following, which effectively says revert the changes from b878a06 to the HEAD at tip of the branch. We also specify `--no-edit` so that the commit messages are automatically applied. This is for convenience now, though normally we would write these commit messages manually using the usual procedure.

```
git revert b878a06^..HEAD --no-edit
```

Removing models.R

```
[master 720ecdb] Revert "Add new lm script named `models.R`"
Date: Sun May 24 22:07:30 2020 +0100
1 file changed, 2 deletions(-)
delete mode 100644 models.R
[master b21c950] Revert "Change function used in the apply functional"
Date: Sun May 24 22:07:30 2020 +0100
1 file changed, 1 insertion(+), 1 deletion(-)
```

From the output of this command, we can see that Git has undone commit b878a06 and then undone commit b878a06. In other words, it is reversing the actions that were applied to produce b878a06, and that then puts it in a state identical to b878a06. Then it reverses the actions that were applied to produce b878a06 so that it is in a state identical to that of 4cb519c.

If we look at the logs in brief, we see the following.

```
git log --pretty=format: '%h %s'
```

```
b21c950 Revert "Change function used in the apply functional"
720ecdb Revert "Add new lm script named `models.R`"
37790d0 Add new lm script named `models.R`
b878a06 Change function used in the apply functional
4cb519c Initialize the repository
```

We see that there are two new commits: b21c950 and 720ecdb. The 720ecdb reverses 37790d0, and then b21c950 reverses b878a06. Notice that the original history has not changed at all: we've added new commits that undo previous ones, but the previous ones are still preserved in the history.

## Branching and merging

Thus far, we have been working with a single branch. By default, this branch is known as the master branch. The following command will tell us about all the branches available in the repository.

```
git branch
```

```
* master
```

In our case, we have just one, `master`, and the `* master` tells that we are on that branch. At any point, we can create a new branch with the `git branch` command. If, for example, we do the following, we create a new branch called `development`.

```
git branch development
```

If we now, re-run `git branch`, we will see the following.

```
git branch
```

```
development
* master
```

This tells us that a new branch, `development`, exists but we are still on `master`. We can switch over to `development` as follows.

```
git checkout development
```

```
Switched to branch 'development'
```

Again, we can re-run `git branch`.

```
git branch
```

```
* development
  master
```

This tells us now that we are on the `development` branch.

There is also a shortcut that allows us to create and switch to a new branch in one command. For example, to create `development` and switch to it immediately, we could do the following.

```
git checkout -B development
```

At this point, we can understand our repository as in the following diagram.

```
o -- o -- o -- o -- o <- tip of master branch
                        \<- tip of development branch
```

Basically, we have two branches but so far both are pointing to the same point in the chain of commits. However, as soon as we stage or commit any edits or new files, these two branches diverge. As an example, let us use `sed` to edit `script.R` by changing the function `sum` to `median`.

```
sed -i 's/sum/median/' script.R
```

Next, we will add this change and commit it using a one-liner commit message.

```
git add script.R
```

```
git commit -m 'Change sum to median in apply functional'
```

```
[development 41640a4] Change sum to median in apply functional
1 file changed, 1 insertion(+), 1 deletion(-)
```

It should be mentioned that we could accomplish both these steps in one step using the following command.

```
git commit -am 'Change sum to median in apply functional'
```

In this case, the `-a` adds any modified files to the staging area.

Let us look at the history.

```
git log --pretty=format: '%h %s'
```

```
41640a4 Change sum to median in apply functional
b21c950 Revert "Change function used in the apply functional"
720ecdb Revert "Add new lm script named `models.R`"
37790d0 Add new lm script named `models.R`
b878a06 Change function used in the apply functional
4cb519c Initialize the repository
```

So far, everything seems like it is behaving as it was before. In other words, it appears as if we have just added another commit to the tip of a chain of commits. That, in fact, is exactly what has happened, but the commit is put on the tip of the `development` branch. The situation can be visualized as follows, where we use `u` to indicate the new commit, i.e. `41640a4`, that we have just created.

```
o -- o -- o -- o -- o <- tip of master branch
                \ -- u <- tip of development branch
```

We can return to the `master` branch as follows.

```
git checkout master
```

Switched to branch 'master'

Again, we can review the logs.

```
git log --pretty=format: '%h %s'
```

```
b21c950 Revert "Change function used in the apply functional"
720ecdb Revert "Add new lm script named `models.R`"
37790d0 Add new lm script named `models.R`
b878a06 Change function used in the apply functional
4cb519c Initialize the repository
```

We see the tip of this branch is `b21c950` and commit `41640a4` is not there.

Now, let us return to the `development` branch and make two changes and then commit each one separately.

```
git checkout development
```

```
sed -i 's/median/sum' script.R
```

```
git commit -am 'Change median to sum in apply functional'
```

```
cat << EOF >> readme.md
```

```
We calculate the number of males and females who survived or not.
EOF
```

```
git commit -am 'Update readme.md'
```

Switched to branch 'development'

```
sed: -e expression #1, char 12: unterminated `s' command
```

```
On branch development
nothing to commit, working tree clean
[development 80e7c2e] Update readme.md
1 file changed, 1 insertion(+)
```

Having added two new commits to the master branch, we have the following situation, where we use `u` to signify the new commits on the development branch.

```
o -- o -- o -- o -- o <- tip of master branch
      \ -- u -- u -- u <- tip of development branch
```

Clearly, `master` and `development` have diverged. We can, however, *merge* these changes on `development` back into `master`. There are many ways of doing this, with heated debates about what best practices should be. For now, we will simply follow the default approach without any commitment as to whether it is the best options of the available alternatives. This default approach is where we `checkout` the `master` and then `merge` in the `development` branch as follows.

```
git checkout master
git merge development
```

```
Switched to branch 'master'
Updating b21c950..80e7c2e
Fast-forward
 readme.md | 1 +
 script.R  | 2 +-
 2 files changed, 2 insertions(+), 1 deletion(-)
```

At this point, all the three changes on the `development` branch are merged into the `master`. The situation now is as follows.

```
o -- o -- o -- o -- o -- u -- u -- u <- tip of master branch
      \ <- tip of development branch
```

Both `development` and `master` again point to the same commit, which is the tip of both branches. We can now, if we wish, delete the master branch. Given that the changes from the master branch have been merged in, we will lose no information. We can delete the branch as follows.

```
git branch -D development
```

Deleted branch development (was 80e7c2e).

If we now run the `git branch` command, we will see that again we have just one branch, namely `master`.

```
git branch
```

```
* master
```

## Using remote repositories

Thus far, we have been using Git locally on one computer, and there has been only one user. However, one of the key reasons for using Git, or any other VCS system, is for sharing and collaborating on projects. For this, we need to use *remote* repositories. It is both inexpensive, or possibly even free, and not technically difficult to host your own private Git server, which can then be used for either private team work or for sharing projects with the public. The required software, including the operating system of the server, i.e. Linux, is open-source and free, and so the only expense is hiring a server from a hosting company<sup>4</sup> However, we will not consider this option further here. Instead, we will consider special purpose Git hosting sites, particularly GitHub. GitHub is extremely popular, with around over 35 million users and 100 million repositories as of November, 2019.

---

<sup>4</sup>From companies such as <https://www.digitalocean.com/>, this may be between \$5-\$10 USD per month.

```
knitr::include_graphics(file.path(chapter_wd, "github_screenshot.png"))
```

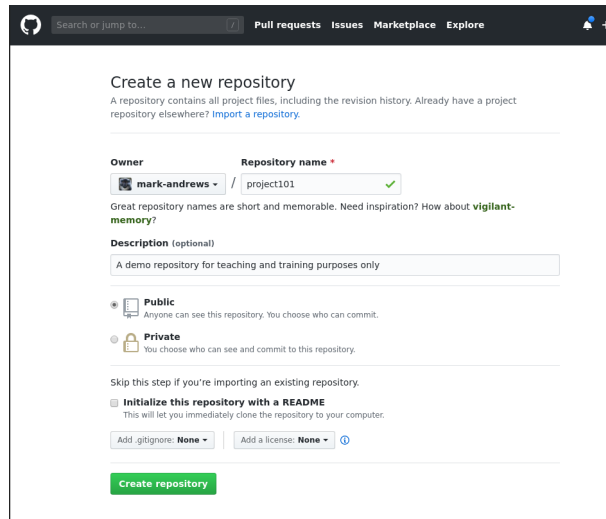


Figure 5: A screenshot of a GitHub page for creating a new repository.

To share your Git repository, the first step is to create a new repository on GitHub. This assumes that you already have a GitHub account. GitHub accounts are free unless you want only privately viewable repositories, which we will not consider here, and you may create unlimited number of publicly viewable repositories using a free account. Assuming you have a GitHub account and have logged, we then can browse to the following url.

<https://github.com/new>

This will bring up a web page that will allow us to create a new empty repository. See Figure 5 for a screenshot of this page. When asked for the repository name, you can use any name. In Figure 5, we use `project101`, but it is not necessary to use the same name for the remote and local repository. The contents of the repository will define it, not its name. You can then add a description. This is just something for the GitHub listing itself. This description will not be part of the repository in general. Then, assuming you've already created a local repository that you wish to push to GitHub, leave everything else on this page at its default. Specifically, do *not* add a `readme`, or a licence, or a `.gitignore` file. After we click the "Create repository", we are then brought to a page that provides some Git commands for different situations. We already have an existing repository, so we want the code listed under the heading *...or push an existing repository from the command line*, which is

```
git remote add origin git@github.com:mark-andrews/project101.git
git push -u origin master
```

There are two important commands here, and so we will look at them individually. The first command adds `github.com:mark-andrews/project101.git` as a remote host of the repository we have created. In this repository, this remote repository will be named `origin`. The name `origin` is the default name for remote repositories from which other repositories are *cloned*. However, we do not have to use this name, and in fact another name, such as `github`, might be more useful, especially if we have multiple remote repositories. The second command *pushes* the contents of the master branch of local repository to the remote repository named `origin`. When you run this, you will be asked for your GitHub password, unless you have set up passwordless GitHub authentication using ssh keys, which is a very convenient feature when frequently using GitHub.

If we run these two commands, if we then browse in a web browser to <https://github.com/mark-andrews/project101> we will see that our repository and all its history and other vital information is being now hosted there.

## Cloning remotes

If you create a public GitHub repository, as we did above, anyone can now *clone* your repository. For example, in a MacOS or Linux terminal or the Windows Git Bash shell, anyone can type the following command and then clone `project101`.

```
git clone git@github.com:mark-andrews/project101.git
```

Alternatively, they could use this version of the clone command.

```
git clone https://github.com/mark-andrews/project101.git
```

The difference between these two versions is simply the internet protocol that is being used. Having cloned `project101`, they will have access to everything that you pushed to GitHub, i.e. the master branch and all the history and other vital features of the repository. They will be able to do everything you can do with the repository: view the logs, make changes, commit changes, undo changes, roll back history, etc. All the logs, commit messages, commit hashes etc will be identical in the clone as in the original. However, as is probably obvious, whatever actions they take, they will not be able to affect your local repository in any way. Moreover, they will not be able to affect your remote repository on GitHub either. They would only be able to **push** to your GitHub repository if they had a GitHub account and you explicitly gave them **push** permission.

## References

- Chacon, Scott, and Ben Straub. 2014. *Pro Git*. 2nd ed. Apress.
- Fecher, Benedikt, Sascha Friesike, and Marcel Hebing. 2015. “What Drives Academic Data Sharing?” *PloS One* 10 (2): e0118053.
- Gorgolewski, Krzysztof J, and Russell A Poldrack. 2016. “A Practical Guide for Improving Transparency and Reproducibility in Neuroimaging Research.” *PLoS Biology* 14 (7): e1002506.
- Grätzer, George. 2016. *Math into LaTeX*. 5th ed. Springer Science & Business Media.
- Gruber, John. 2004. “Markdown.” <https://daringfireball.net/projects/markdown/>.
- Hern, Alex. 2013. “Is Excel the Most Dangerous Piece of Software in the World?” *New Statesman*. <https://www.newstatesman.com/technology/2013/02/excel-most-dangerous-piece-software-world>.
- Houtkoop, Bobby Lee, Chris Chambers, Malcolm Macleod, Dorothy VM Bishop, Thomas E Nichols, and Eric-Jan Wagenmakers. 2018. “Data Sharing in Psychology: A Survey on Barriers and Preconditions.” *Advances in Methods and Practices in Psychological Science* 1 (1): 70–85.
- Ioannidis, John PA. 2015. “How to Make More Published Research True.” *Revista Cubana de Información En Ciencias de La Salud (ACIMED)* 26 (2): 187–200.
- Iqbal, Shareen A, Joshua D Wallach, Muin J Khoury, Sheri D Schully, and John PA Ioannidis. 2016. “Reproducible Research Practices and Transparency Across the Biomedical Literature.” *PLoS Biology* 14 (1): e1002333.
- Knuth, Donald Ervin. 1984. “Literate Programming.” *The Computer Journal* 27 (2): 97–111.
- MacFarlane, John. 2006. “Pandoc: A Universal Document Converter.” <https://pandoc.org/>.
- Merton, Robert K. 1973. *The Sociology of Science: Theoretical and Empirical Investigations*. University of Chicago press.
- Munafò, Marcus R, Brian A Nosek, Dorothy VM Bishop, Katherine S Button, Christopher D Chambers, Nathalie Percie Du Sert, Uri Simonsohn, Eric-Jan Wagenmakers, Jennifer J Ware, and John PA Ioannidis. 2017. “A Manifesto for Reproducible Science.” *Nature Human Behaviour* 1 (1): 0021.



- Nosek, Brian A, George Alter, George C Banks, Denny Borsboom, Sara D Bowman, Steven J Breckler, Stuart Buck, et al. 2015. “Promoting an Open Research Culture.” *Science* 348 (6242): 1422–5.
- Shamir, Lior, John F Wallin, Alice Allen, Bruce Berriman, Peter Teuben, Robert J Nemiroff, Jessica Mink, Robert J Hanisch, and Kimberly DuPrie. 2013. “Practices in Source Code Sharing in Astrophysics.” *Astronomy and Computing* 1: 54–58.
- Stodden, Victoria, Peixuan Guo, and Zhaokun Ma. 2013. “Toward Reproducible Computational Research: An Empirical Analysis of Data and Code Policy Adoption by Journals.” *PloS One* 8 (6): e67111.
- Tenopir, Carol, Suzie Allard, Kimberly Douglass, Arsev Umur Aydinoglu, Lei Wu, Eleanor Read, Maribeth Manoff, and Mike Frame. 2011. “Data Sharing by Scientists: Practices and Perceptions.” *PloS One* 6 (6): e21101.
- Xie, Yihui. 2017. *Dynamic Documents with R and Knitr*. Chapman; Hall/CRC.